

APPSEC-10: Dashboards, Reporting and Governance

Series: APPSEC — Application Security | **Notebook:** 10 of 10 | **Created:** June 2026
| **Last Updated:** 09/28/2026

Overview

AppSec produces signal continuously — dashboards and governance turn that signal into decisions. This closing notebook covers the executive-facing dashboard composition, the cadence of the governance review, and the FinOps angle (RVA and RAP are DPS-billed per monitored GiB-hour, so cost follows where they are enabled).

This is the final notebook in the series. The other nine cover the data sources; this one covers what to do with them at the leadership and program level.



Table of Contents

1. [1. Open Vulnerabilities by DSS Level — the Trend](#)
 2. [2. Open Problems by Severity](#)
 3. [3. MTTR by Team](#)
 4. [4. Compliance Framework Coverage](#)
 5. [5. Governance Cadence](#)
 6. [6. DPS Cost Awareness](#)
 7. [7. Series Wrap](#)
 8. [References](#)
-

Prerequisites

Requirement	Details
Dynatrace Environment	Gen3 SaaS with Grail; AppSec entitlement enabled
OneAgent	Full-Stack mode (or code-module attached) on monitored hosts
Read access	At minimum <code>environment:roles:view-security-problems</code> and <code>storage:security.events:read</code> — see APPSEC-09 for the full model
Background	APPSEC-01 (fundamentals + three-pillar framing)

1. Open Vulnerabilities by DSS Level — the Trend

DSS is a per-vulnerability score (APPSEC-01 § 4), so there is no single tenant-wide "DSS" to chart. The board-facing metric is the **count of open vulnerabilities at each DSS level**, trended over a 90-day window — § 2's query, run on a schedule. In community practice the trend is read directionally:

- **More open vulnerabilities at CRITICAL / HIGH** = posture deteriorating. Investigate which applications or libraries added them.

- **Flat** = posture stable. Could mean no new risk, or could mean stalled remediation — pair with backlog burn-rate from APPSEC-08 to disambiguate.
- **Fewer open vulnerabilities at CRITICAL / HIGH** = remediation is outpacing new findings.

Do **not** compare these counts across tenants or business units without normalizing for estate size and monitoring-mode coverage (per APPSEC-01 § 3–4) — Infrastructure-mode hosts keep DSS at the CVSS base score, which shifts the level mix. Use within-tenant deltas.

Sources: [Vulnerabilities concepts \(DT docs\)](#) — *"This scoring system forms the foundation for the Dynatrace Security Score (DSS), which adds environmental context to help prioritize remediation."*; [Application Security \(DT docs\)](#) — *"the DSS will be the same as the CVSS base score"* in Infrastructure Monitoring deployments (both re-read 09/18/2026). **Derived:** the cross-tenant caution follows from DSS falling back to the CVSS base score in Infrastructure Monitoring.

2. Open Problems by Severity

The underlying view: how many vulnerabilities are open right now, by DSS level? The query keeps the latest vulnerability-level snapshot per vulnerability, so each open vulnerability counts once however often RVA re-reports it.

Severity	Treatment
Critical	Daily review; SLA 7 days
High	Weekly review; SLA 30 days
Medium	Sprint review; SLA 90 days
Low	Quarterly cleanup

A healthy mix has Critical near zero, High in single digits, Medium and Low forming the long tail. A growing Critical count is the strongest early-warning signal.

```
// Current open vulnerabilities by risk level (latest snapshot per
vulnerability)
fetch security.events, from:-7d
| filter event.provider == "Dynatrace"
| filter event.category == "VULNERABILITY_MANAGEMENT"
| filter event.type == "VULNERABILITY_STATE_REPORT_EVENT"
| filter event.level == "VULNERABILITY"
| dedup {vulnerability.display_id}, sort:{timestamp desc}
| filter vulnerability.resolution.status == "OPEN"
| filter vulnerability.mute.status == "NOT_MUTED"
| summarize open = count(), by:{vulnerability.risk.level}
| sort open desc
```

Validation note: validated for syntax and field names on 09/18/2026 and executes cleanly; the validation tenant has no RVA data, so it returned 0 rows there.

Sources: [Vulnerability events \(DT semantic dictionary\)](#) — "A vulnerability state event is a periodic snapshot, emitted by Dynatrace Runtime Vulnerability Analytics (RVA), of a vulnerability's current status and risk aggregated across every entity it affects." (re-read 09/18/2026). **Dictionary:**

`vulnerability.risk.level` (stable), `vulnerability.resolution.status` (stable), `vulnerability.mute.status` (stable), `vulnerability.display_id` (stable), read 09/18/2026. **Softened:** the SLA day counts (7/30/90) are common but should be set per your governance regime, not adopted blindly.

3. MTTR by Team

In SRE and incident-management practice, Mean Time To Remediate by team is used as the fairness metric — the AppSec docs do not prescribe it. It tells you which teams are keeping up and which need help (engineering capacity, blocked dependencies, prioritization conflict).

In community practice, accountability dashboards group by team or by namespace owner — not by individual developer, which tends to create the wrong incentives and surface noise.

4. Compliance Framework Coverage

For organizations subject to specific compliance regimes (PCI, HIPAA, SOC 2, ISO 27001), in community practice the dashboard view that matters is *percent of controls in compliance* by framework, not total finding counts.

Pick the primary framework (one) for governance reporting and use the others as secondary views — community guidance is that cross-framework rollups mislead because the same finding lands in multiple frameworks; verify against your audit regime.

Sources: [Security Posture Management \(DT docs\)](#) — "*Automated assessments against supported compliance standards*" (re-read 09/28/2026).

5. Governance Cadence

In community practice, a workable cadence for most organizations looks like this — adapt it to your existing security-governance rhythm:

Cadence	Audience	Content
Daily	Security on-call	New Critical vulnerabilities, RAP detection spikes
Weekly	Security engineering	Open backlog by team, MTTR trend, burn-rate
Monthly	AppSec program lead + AppDev leadership	Open-vulnerability trend by DSS level, compliance coverage, IAM access review
Quarterly	CISO + exec staff	90-day trend of open vulnerabilities by DSS level, compliance posture by framework, capacity asks

The artifacts above (open-vulnerability trend by DSS level, severity mix, MTTR, compliance) feed all four cadences — what differs is the aggregation level and the audience-appropriate framing.

6. DPS Cost Awareness

Runtime Vulnerability Analytics and Runtime Application Protection are DPS-billed in **GiB-hours** — per monitored host or container, based on its memory — not per `security.events` record. Cost is therefore driven by *which hosts have RVA and RAP enabled*, not by how many findings or detections they produce. Three FinOps practices:

1. **Scope enablement, not events.** Use monitoring rules to enable RVA and RAP where the risk justifies it; dropping or sampling detection records saves nothing and discards attack evidence.
2. **Track consumption per capability and host** — in community practice AppSec is tracked as one tenant-wide consumer among many; FINOPS-01 § 5 covers the host-based capability queries, and FINOPS-02 the forecasting model.
3. **Watch detection volume as a signal-quality metric, not a cost metric.** In community practice, a spike in RAP detections with no matching incident is read as rules generating false positives — tune them (APPSEC-04 § 3) for the analysts' sake.

Don't sacrifice security signal for cost savings without an explicit risk acceptance. But don't ignore cost either — it's a budget reality.

Sources: [Runtime Vulnerability Analytics \(DPS\) \(DT docs\)](#) — "*The unit of measure for Runtime Vulnerability Analytics is the GiB-hour*"; [Runtime Application Protection \(DPS\) \(DT docs\)](#) — "*The unit of measure for Runtime Application Protection is a GiB hour*" (both re-read 09/18/2026); [Runtime Vulnerability Analytics \(DT docs\)](#) — "*You can also define custom monitoring rules based on certain criteria.*"; [Runtime Application Protection \(DT docs\)](#) — "*If you define custom monitoring rules based on certain process groups or vulnerability types, the custom rules override the global attack control*" (both re-read 09/28/2026). **Derived:** practice 1 follows from the GiB-hour unit plus rule-based enablement.

7. Series Wrap

The series in one paragraph: AppSec in Gen3 SaaS is **three pillars** (RVA + RAP + SPM) over **one Grail data plane** (`security.events` + `vulnerability-service`), with **dual-surface IAM** (Grail + environment roles), consumed via **dashboards / workflows / Davis CoPilot**, and governed at the **monthly + quarterly** cadence with **open vulnerabilities by DSS level** as the headline metric. The OneAgent code module is the load-bearing dependency under all of it.

Where to go from here:

- Open APPSEC-01 again with the rest of the series fresh and confirm the mental map.
- Stand up the persona policies from APPSEC-09 before granting broad access.
- Wire one end-to-end workflow per APPSEC-08 to prove the loop closes.
- Set the governance cadence and the first review date for the open-vulnerability trend.

References

Source	Coverage
Application Security (DT docs)	DSS + posture framing
Vulnerabilities concepts (DT docs)	DSS as a per-vulnerability score
Vulnerability events (DT semantic dictionary)	Fields used in § 2
Runtime Vulnerability Analytics (DPS) (DT docs)	RVA billing unit
Runtime Application Protection (DPS) (DT docs)	RAP billing unit

 **DISCLAIMER:** This information was AI generated and is provided "as-is" without warranty. It was produced as an independent, community-driven project and **not supported by Dynatrace**. Always refer to official [Dynatrace documentation](#) for the most current information.