

DASH-99: Best Practice Summary

Series: DASH — Dashboard Design & Building | **Notebook:** 99 | **Created:** March 2026 |
Last Updated: 09/28/2026

Overview

This notebook consolidates every actionable best practice from the DASH series (DASH-01 through DASH-07) into a single definitive reference. Each practice specifies exactly what to set, with no ambiguity. Use this as a checklist when building, auditing, or reviewing Dynatrace dashboards.

Table of Contents

1. [Dashboard Architecture](#)
 2. [Design and Layout](#)
 3. [Tile Configuration](#)
 4. [Executive Tier](#)
 5. [Operations Tier](#)
 6. [Engineering Tier](#)
 7. [DQL Query Standards](#)
 8. [Variables and Filters](#)
 9. [Refresh and Performance](#)
 10. [Sharing and Permissions](#)
 11. [Reporting and Automation](#)
 12. [Dashboard Lifecycle](#)
-

Prerequisites

Requirement	Details
Prior Reading	DASH-01 through DASH-07 for full context
Dynatrace Environment	SaaS or Managed with Grail enabled
Permissions	<code>storage:logs:read</code> , <code>storage:metrics:read</code> , <code>storage:events:read</code> , <code>storage:spans:read</code> , <code>document:documents:write</code>

1. Dashboard Architecture

#	Best Practice	Recommended Setting/Value	Priority	Source
1	Use the three-tier hierarchy	Create exactly three tiers: Executive, Operations, Engineering	Critical	DASH-02
2	One purpose per dashboard	Each dashboard answers one audience's questions; never build an "everything" dashboard	Critical	DASH-01
3	Connect tiers with drill-down links	Add markdown tiles with dashboard links so users navigate Executive > Operations > Engineering	Recommended	DASH-02
4	Use dashboards for monitoring, notebooks for investigation	Dashboards = continuous monitoring, team visibility, stakeholder reporting. Notebooks = ad-hoc exploration, incident investigation, training	Critical	DASH-01
5	Organize tiles into sections	Group related tiles into named sections (e.g., "Service Health", "Infrastructure"); sections can be collapsed	Recommended	DASH-01

2. Design and Layout

#	Best Practice	Recommended Setting/Value	Priority	Source
6	Place the most critical KPI in the top-left	Top-left tile is the first thing the eye hits; put your primary health indicator there	Critical	DASH-01
7	Follow Z-pattern reading order	Arrange tiles left-to-right, top-to-bottom so the dashboard reads naturally	Recommended	DASH-01
8	Limit tile count to 8-12 per dashboard	More than 15 tiles creates cognitive overload	Critical	DASH-01
9	Use descriptive tile titles	Title format: Metric Name (Unit) – Scope . Example: "Error Rate (%) — Checkout Service". Never use "Query 1"	Critical	DASH-01
10	Keep all tiles on the same time range	Use the dashboard time selector for all tiles; only use fixed ranges when there is a specific analytical reason	Recommended	DASH-01
11	Use color intentionally	Green = healthy, Yellow = watch, Red = act. No decorative colors	Recommended	DASH-03
12	No decorative elements	No gratuitous pie charts, no 3D effects, no logos	Optional	DASH-03
13	Executive layout: headline-trend-detail	Top row: 3 single-value KPIs. Middle row: trend line chart. Bottom row: detail table	Recommended	DASH-03

3. Tile Configuration

#	Best Practice	Recommended Setting/Value	Priority	Source
14	Every tile must trigger an action	If a tile does not lead to a decision (investigate, escalate, scale, correlate), remove it	Critical	DASH-01

#	Best Practice	Recommended Setting/Value	Priority	Source
15	Match tile type to data shape	Single value: KPIs. Line chart: trends over time. Bar chart: categorical comparisons. Table: detailed data / top-N. Honeycomb: entity health grid	Recommended	DASH-01
16	Use single-value tiles for executive KPIs	Availability %, MTTR, active problem count, error budget remaining	Critical	DASH-03
17	Use line charts for time trends	Problem count over 7d, error rate over 2h, CPU usage over 1h	Recommended	DASH-02
18	Use table tiles for engineering deep-dives	Engineers need exact values (endpoint p95, DB query time, trace IDs), not visual approximations	Recommended	DASH-05

4. Executive Tier

#	Best Practice	Recommended Setting/Value	Priority	Source
19	Limit to 4-6 tiles maximum	Single-value and trend tiles only	Critical	DASH-02, DASH-03
20	Select KPIs using the SMART filter	Each KPI must be Specific, Measurable, Actionable, Relevant, Time-bound	Critical	DASH-03
21	Start with exactly these 5 KPIs	Availability %, MTTR, Active Problem Count, Problem Trend (7d), Overall Error Rate	Critical	DASH-03
22	Set availability thresholds	Green: >99.9%, Yellow: >99.5%, Red: <99.5%	Critical	DASH-03
23	Set MTTR thresholds	Green: <1 hour, Yellow: <4 hours, Red: >4 hours	Critical	DASH-03
24	Set active problem thresholds	Green: 0, Yellow: 1-3, Red: >3	Recommended	DASH-03

#	Best Practice	Recommended Setting/Value	Priority	Source
25	Set error rate thresholds	Green: <1%, Yellow: <5%, Red: >5%	Recommended	DASH-03
26	Track error budget for SLA targets	99.9% SLA = 43.2 min/month budget. 99.5% = 3.6 hours. 99.0% = 7.2 hours	Recommended	DASH-03
27	Use 7d-30d time ranges	Executive dashboards show weekly or monthly trends, not minute-by-minute data	Critical	DASH-02
28	No technical jargon	"Availability" not "HTTP 200 ratio". "Resolution Time" not "MTTR" if audience is non-technical	Recommended	DASH-02
29	Exclude duplicate and frequent detected problems	Always filter: <code>dt.davis.is_duplicate == false</code> and <code>dt.davis.is_frequent_event == false</code>	Critical	DASH-03
30	Dashboard must be understandable in <30 seconds	If it takes longer, push detail down to operations tier	Critical	DASH-02

5. Operations Tier

#	Best Practice	Recommended Setting/Value	Priority	Source
31	Use 8-12 tiles with mixed chart types	Line charts for trends, bar charts for comparisons, tables for top-N, single values for status	Critical	DASH-02, DASH-04
32	Use 1h-2h time windows for real-time tiles	Short windows surface real-time spikes; 24h hides them	Critical	DASH-02, DASH-04
33	Show p50, p90, p95 response time lines	Three percentile lines on a single timeseries chart at 5-minute intervals	Recommended	DASH-04

#	Best Practice	Recommended Setting/Value	Priority	Source
34	Include active problems table	Show display_id, event.name, event.category, duration_min. Sort by duration descending	Critical	DASH-04
35	Track log volume by severity level	Stacked area chart: <code>makeTimeseries count(), interval:5m, by:{loglevel} .</code> Filter out <code>loglevel == "NONE"</code>	Recommended	DASH-04
36	Detect log volume spikes	Compare current hour log count to same hour yesterday using <code>append</code>	Recommended	DASH-04
37	Include deployment event table	Filter <code>event.type == "CUSTOM_DEPLOYMENT"</code> or <code>(event.kind == "SDLC_EVENT" and event.category == "task" and event.type == "deployment")</code> — classic and SDLC deployment events. Show last 6h	Critical	DASH-04
38	Add infrastructure tiles for CPU and memory	<code>timeseries avg(dt.host.cpu.usage)</code> and <code>timeseries avg(dt.host.memory.usage)</code> by host	Recommended	DASH-04
39	Show concurrently open problems over time	Use <code>spread:timeframe(from:event.start, to:coalesce(event.end, now()))</code> in <code>makeTimeseries</code>	Recommended	DASH-04
40	Avoid spaghetti line charts	Use top-N filtering or variables instead of plotting all services on one chart	Critical	DASH-04

6. Engineering Tier

#	Best Practice	Recommended Setting/Value	Priority	Source
41	Use 10-15 tiles with high data density	Tables, detailed charts, multiple variables	Recommended	DASH-02, DASH-05
42	Use 15m-1h time ranges for investigation	Longer ranges dilute the signal; shorter ranges focus on the incident window	Critical	DASH-05
43	Include service throughput table	Show request_count, avg_ms, p50_ms, p95_ms, p99_ms, error_rate_pct per service	Critical	DASH-05
44	Include endpoint performance table	Show requests, avg_ms, p95_ms, error_rate by http.route and http.request.method	Critical	DASH-05
45	Include database performance table	Show avg_ms, p95_ms, query_count, error_rate by db.system and db.namespace	Recommended	DASH-05
46	Always filter <code>isNotNull(db.system)</code> on DB spans	Prevents null groups in database analysis	Critical	DASH-05
47	Include latency distribution histogram	Bucket spans into <100ms, 100-500ms, 500ms-1s, 1s-5s, >5s categories	Recommended	DASH-05
48	Provide before/after deployment comparison	Use <code>append</code> to compare 2h after vs 2h before for avg_ms, p95_ms, error_rate	Recommended	DASH-05
49	Include trace.id in error span tables	Engineers need clickable trace IDs for distributed trace investigation	Critical	DASH-05

#	Best Practice	Recommended Setting/Value	Priority	Source
50	Add markdown tiles documenting each section	Explain what the section shows and how to interpret it	Recommended	DASH-05
51	Use variables extensively	Service, namespace, time range selectors so engineers can focus on their area	Critical	DASH-05

7. DQL Query Standards

#	Best Practice	Recommended Setting/Value	Priority	Source
52	Always specify a time range on fetch	<code>fetch logs, from: -1h</code> . Never use bare <code>fetch logs</code>	Critical	DASH-01, all
53	Filter before aggregating	Place <code>filter</code> immediately after <code>fetch</code> , never after <code>sort</code> or <code>summarize</code>	Critical	DASH-01
54	Alias all aggregation results	<code>summarize c = count()</code> not <code>summarize count()</code> . Required for sort/fieldsAdd references	Critical	All
55	Convert span duration to milliseconds	<code>duration / 1ms</code> for display. Duration is stored in nanoseconds	Critical	DASH-04, DASH-05
56	Convert detected problem duration to hours	<code>resolved_problem_duration / 1h</code> . Duration is in nanoseconds	Critical	DASH-03
57	Use <code>arrayAvg()</code> on timeseries results before filter/sort	Timeseries returns arrays, not scalars. Use <code>fieldsAdd val = arrayAvg(ts)</code> then <code>filter val > x</code>	Critical	DASH-04

#	Best Practice	Recommended Setting/Value	Priority	Source
58	Use <code>countIf()</code> for conditional counts	<code>errors = countIf(span.status_code == "error")</code> instead of separate filter + count	Recommended	DASH-03
59	Calculate error rate as percentage	<code>100.0 * errors / total</code> . Always multiply by 100.0 (float), not 100 (int)	Recommended	DASH-04
60	Filter out duplicate detected problems	Always include <code>filter dt.davis.is_duplicate == false</code> in problem queries	Critical	DASH-03
61	Use <code>round()</code> with named <code>decimals:</code> parameter	<code>round(value, decimals: 2)</code> not <code>round(value, 2)</code>	Critical	All
62	Prototype all queries in a notebook first	Validate DQL in a notebook before transferring to dashboard tiles	Recommended	DASH-01
63	Use <code>span.status_code</code> for span failure — never <code>otel.status_code</code>	<code>otel.status_code</code> and <code>otel.status_message</code> do not exist (no row in <code>dt.semantic_dictionary.fields</code>). The real fields are <code>span.status_code (stable)</code> and <code>span.status_message (experimental)</code>	Critical	DASH-03, DASH-04, DASH-05
64	Match <code>span.status_code</code> in lowercase	<code>== "error"</code> , never <code>== "ERROR"</code> . The OTel SDK constant is uppercase; the Grail value is not. On a live tenant <code>"ERROR"</code> matched 0 spans against 18,630 for <code>"error"</code> (08/11/2026)	Critical	DASH-03, DASH-04, DASH-05
65	Derive span successes as <code>total - errors</code>	<code>span.status_code</code> is null on successful spans — <code>"ok"</code> is written vanishingly rarely — so <code>countIf(span.status_code != "error")</code> counts almost nothing and an availability tile reads ~0% on a healthy estate	Critical	DASH-03

Rules 63–65 are one bug, caught three ways. Each of the three mistakes returns a *plausible number* rather than an error, so nothing in the tile, the query, or the dashboard tells you it is wrong. The order in which they bite matters: fixing only the field name (63) turns a hard failure into a silent zero (64), and fixing the name and the case but not the null-handling (65) turns a correct-looking rename into a **0% availability** executive tile on a service that is actually running at 99.6%. Verified end to end against a live tenant, 08/11/2026.

8. Variables and Filters

#	Best Practice	Recommended Setting/Value	Priority	Source
66	Use DQL variables for services and hosts	Variable type: DQL, with a query over <code>smartscapeNodes</code> or <code>dt.entity.service</code> / <code>dt.entity.host</code> . The Dashboards app has no entity-selector type	Critical	DASH-06
67	Use List variables for fixed sets: namespaces, log levels, environments	Values discovered from actual data via DQL; switch to a DQL variable when the set changes often	Recommended	DASH-06
68	Reference the variable key, not its display name	DQL uses the key, derived from the name with every non-alphanumeric character replaced by <code>_</code> . Keep keys lowercase with underscores: <code>\$k8s_namespace</code> , <code>\$service</code> , <code>\$environment</code>	Critical	DASH-06
69	Reference variables in filter clauses	<code>filter dt.entity.host == \$host</code> for DQL entity variables; <code>filter k8s.namespace.name == \$namespace</code> for List or Free-text; <code>filter in(loglevel, array(\$log_levels))</code> for multi-select	Critical	DASH-06
70	Provide an "All" option on every variable	Let users remove the filter to see aggregate data	Recommended	DASH-06

#	Best Practice	Recommended Setting/Value	Priority	Source
71	Use the same variable name across all related tiles	Ensures filter propagation: changing the dropdown updates every tile that references it	Critical	DASH-06
72	Chain variables for hierarchical filtering	Second variable's query-based options reference the first variable (e.g., cluster > namespace)	Optional	DASH-06
73	Build Golden Signals template dashboards	One variable-driven template with Latency, Traffic, Errors, Saturation sections. Works for any service	Recommended	DASH-06
74	Test variables with extreme values	Select the busiest and quietest entities to verify the layout renders correctly under load	Recommended	DASH-06

9. Refresh and Performance

#	Best Practice	Recommended Setting/Value	Priority	Source
75	Executive wall screen refresh	Set to 10-15 minutes	Critical	DASH-02
76	Operations NOC wall refresh	Set to 1-2 minutes	Critical	DASH-02
77	Operations on-call refresh	Set to 2-5 minutes	Recommended	DASH-02
78	Engineering investigation refresh	Set to manual (on-demand)	Recommended	DASH-02
79	Executive weekly review refresh	Set to manual	Optional	DASH-02
80	Never set refresh below 1 minute	Aggressive refresh on complex queries causes unnecessary load. Start at 5 min, decrease only if critical	Critical	DASH-02

#	Best Practice	Recommended Setting/Value	Priority	Source
81	Use <code>limit</code> on all top-N queries	Always append <code> limit N</code> (10-20) to prevent unbounded result sets in tiles	Critical	DASH-04, DASH-05
82	Use <code>filterOut</code> instead of <code>not</code> for negation filters	<code>filterOut loglevel == "NONE"</code> is faster than <code>filter not loglevel == "NONE"</code>	Recommended	DASH-04

10. Sharing and Permissions

#	Best Practice	Recommended Setting/Value	Priority	Source
83	Executive dashboards: platform team owns, leadership views	Owner: platform team lead. Editors: platform team. Viewers: all managers + leadership	Critical	DASH-07
84	Operations dashboards: SRE team owns	Owner: SRE team. Editors: SRE + on-call. Viewers: operations group	Critical	DASH-07
85	Engineering dashboards: service team owns	Owner: service team lead. Editors: service team. Viewers: engineering org	Recommended	DASH-07
86	Limit editor access	Encourage clone-and-customize over editing shared originals	Critical	DASH-07
87	Use naming convention for all dashboards	Format: <code>[Tier] - [Team/Service] - [Purpose]</code> . Example: "Ops - Checkout - Service Health"	Critical	DASH-07
88	Tag dashboards with tier, team, and service	Enables filtering and discovery in the dashboard library	Recommended	DASH-07

#	Best Practice	Recommended Setting/Value	Priority	Source
89	Maintain a dashboard ownership registry	Document who owns each dashboard for accountability	Recommended	DASH-07
90	Archive dashboards not used in 90 days	Proxy: no tile queries from the Dashboards app in 90 days (<code>QUERY_EXECUTION_EVENT</code> in <code>dt.system.events</code>) — document reads are not audited. Regular cleanup prevents library bloat	Recommended	DASH-07

11. Reporting and Automation

#	Best Practice	Recommended Setting/Value	Priority	Source
91	Automate executive summary via Workflow	Schedule: weekly, Monday 8 AM. Recipients: leadership, management	Recommended	DASH-07
92	Automate operations daily report via Workflow	Schedule: daily, 7 AM. Recipients: SRE team, on-call	Recommended	DASH-07
93	Automate SLA compliance report via Workflow	Schedule: monthly, 1st of month. Recipients: account management, leadership	Recommended	DASH-07
94	Automate cost/volume tracking report	Schedule: weekly. Recipients: platform team, finance	Optional	DASH-07

#	Best Practice	Recommended Setting/Value	Priority	Source
95	Only report services with error rate >1% in daily ops report	Filter <code>error_rate_pct > 1.0</code> to surface only degraded services	Recommended	DASH-07
96	Export dashboard JSON for backup	Use Documents API: <code>GET /platform/document/v1/documents/<id></code> returns metadata and content as one multipart response — store the content part in version control	Recommended	DASH-07

12. Dashboard Lifecycle

#	Best Practice	Recommended Setting/Value	Priority	Source
97	Follow the 6-step creation workflow	1) Define audience/purpose, 2) Identify 5-10 key metrics, 3) Prototype in notebook, 4) Build dashboard, 5) Add variables, 6) Share and iterate	Critical	DASH-01
98	Manage dashboards as code	Use Monaco CLI or Terraform for multi-environment deployment, version control, and peer review	Recommended	DASH-07
99	Validate every dashboard payload before merging it	Deploy to a non-production tenant, open the dashboard in the Dashboards app, confirm zero validation warnings. API acceptance (<code>2xx</code>) is not render confirmation. Since SaaS 1.344 (stricter rules in SaaS 1.346) a dashboard that fails validation does not load at all — confirm which version your tenant runs — and API- or AI-authored dashboards are the named affected population	Critical	DASH-07

#	Best Practice	Recommended Setting/Value	Priority	Source
100	Organize dashboard repo by tier	Structure: dashboards/executive/ , dashboards/operations/ , dashboards/engineering/ , dashboards/templates/	Recommended	DASH-07
101	Deploy dashboard changes through PR review	Export > modify in feature branch > PR review > deploy to staging > validate schema/render > review data fidelity > deploy to production	Recommended	DASH-07
102	Maintain a query registry	Track which DQL queries power which dashboard tiles. Update when DQL syntax or data sources change	Optional	DASH-07
103	A dashboard is never "done"	Share with target audience, collect feedback, and refine continuously. Evolve with team needs	Recommended	DASH-01
104	Validate dashboard data matches manual reports	Executives lose trust if dashboard numbers differ from status meeting reports. Cross-check availability and MTTR calculations	Critical	DASH-03

Summary

This notebook contains **104 best practices** extracted from the DASH series, organized into 12 categories:

Category	Practices	Critical Count
Dashboard Architecture	1-5	3
Design and Layout	6-13	3
Tile Configuration	14-18	2
Executive Tier	19-30	8
Operations Tier	31-40	5

Category	Practices	Critical Count
Engineering Tier	41-51	6
DQL Query Standards	52-65	11
Variables and Filters	66-74	4
Refresh and Performance	75-82	4
Sharing and Permissions	83-90	4
Reporting and Automation	91-96	0
Dashboard Lifecycle	97-104	3

Priority breakdown: 53 Critical, 46 Recommended, 5 Optional. Start with Critical practices first, then layer in Recommended practices for a mature dashboard strategy.

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.