

FAQ-02: Tagging — Sources, Standards, and Strategy

Series: FAQ — Frequently Asked Questions | **Reference:** 02 — Tagging Sources, Standards, and Strategy | **Created:** May 2026 | **Last Updated:** 09/28/2026

Overview

Tagging is one of the most consequential decisions in a Dynatrace tenant — it drives ownership, alerting, IAM scope, automation targeting, cost attribution, and dashboard filtering. It is also one of the most commonly under-designed: teams accumulate tags from multiple sources without a coherent strategy, end up with inconsistent values across clouds, and discover the gaps only when an audit, a cost-allocation report, or an incident-routing rule fails.

This FAQ entry consolidates the four tag sources that flow into Dynatrace (OneAgent, Kubernetes, cloud providers, and legacy auto-tagging rules), the difference between **primary tags** / **primary fields** and ordinary tags, the standards that make a tagging strategy work across an estate, and the strategy themes that turn tag sprawl into a managed asset.

If you read only one section, read **§5 (Standards)** and **§6 (Strategy)** — together they answer most "how should we tag?" questions.

Table of Contents

1. [What "Tags" Mean in Dynatrace](#)
 2. [Primary Tags vs Primary Fields vs Custom Tags vs Auto-Tags](#)
 3. [The Four-Source Hierarchy](#)
 - [Release-identity environment variables \(`DT\_RELEASE\_\*` \)](#)
 4. [AWS / Azure / GCP — Per-Cloud Specifics](#)
 5. [Tagging Standards — Taxonomy and Naming](#)
 6. [Strategy Themes](#)
 7. [Anti-Patterns](#)
 8. [Final Recommendation](#)
 9. [Related Resources](#)
-

Prerequisites

Requirement	Details
Audience	Platform team, SRE leads, FinOps stakeholders, security architects, anyone deciding how to model ownership/attribution in Dynatrace
Format	Decision-support document — presents sources, standards, and strategy; no hands-on lab
Related topic series	OPIPE (primary fields at source, OpenPipeline enrichment), CLOUD (per-provider field mappings), IAM (<code>dt.security_context</code> as boundary), AUTOM (config-as-code for any remaining tag rules), K8S (Kubernetes label propagation), ORGNZ (segments and bucket strategy that consume tags)

1. What "Tags" Mean in Dynatrace

"Tags" is an overloaded word. In a Dynatrace tenant, four distinct things commonly get called "tags":

Concept	Where it lives	When it shows up in Dynatrace
Cloud-provider tags	AWS / Azure / GCP resources	Via the Clouds app or legacy CloudWatch / Azure Monitor / GCP integrations; surface as <code>aws.tag.<key></code> , <code>azure.tag.<key></code> , <code>gcp.label.<key></code> plus provider attributes (<code>aws.account.id</code> , <code>azure.resource_group</code> , <code>gcp.project_id</code>)
Kubernetes labels and annotations	K8s manifests / Helm / GitOps	Via DynaKube + OneAgent metadata enrichment; surface as <code>k8s.<resource>.label.<key></code> (e.g., <code>k8s.pod.label.app</code> , <code>k8s.namespace.label.team</code>)
OneAgent host tags / primary tags	Set on the host at install or via <code>oneagentctl --set-host-tag=</code>	Ride on every signal at source — metrics, spans, logs, business events, Smartscape entities — including the sprint-1.337+ primary fields (<code>dt.security_context</code> , <code>dt.cost.costcenter</code> , <code>dt.cost.product</code>)
Auto-tagging rules (<i>legacy</i>)	Settings 2.0 (or legacy Configuration API) — rules that compute a tag at view time from other entity properties	Computed when an entity is viewed/queried; not present at ingest


These are not interchangeable. They differ in **where** they're set, **when** they're computed, **what** they propagate to, and **how** they're consumed downstream:

- **Where set:** changes to OneAgent host tags happen at the host's edge; cloud tags happen in the provider console; K8s labels happen in your manifests; auto-tagging rules happen in the Dynatrace Settings surface.
- **When computed:** the first three are set at ingest (the value rides on every signal); auto-tagging rules are computed at view/query time.
- **What they propagate to:** primary tags propagate to every signal type; auto-tagging applies primarily at the entity level (and not uniformly to logs, spans, business events).
- **How consumed:** Smartscape, dashboards, alerts/thresholds, IAM policies, OpenPipeline routing, segments — different consumers require different sources to be the source of truth.

The implication: **picking the right source for each dimension you tag on is more important than picking the right tag value.** A consistent value carried in the wrong source is harder to fix than a typo.

Sources: [Tags and metadata \(DT docs\)](#) — seven documented categories of tags and metadata; *"Host groups can also be used in tagging rules and for defining management zones".*

2. Primary Tags vs Primary Fields vs Custom Tags vs Auto-Tags



The layering is why tagging is worth doing carefully. A value set at the source is on the record before routing, and from there the same value is what access, alert routing, chargeback and compliance evidence all read: *"Dynatrace enriches all derived signals (service metrics, Davis events, and problems) with the same tags."* It is also why ownership belongs in a tag rather than a dashboard filter — Dynatrace names the use case directly: *"Route alert notifications to the right recipients based on the ownership metadata on every Davis event."*

Spaces — rolling out with SaaS 1.348 (pre-release; staged tenant rollout planned from 09/22/2026).

The release notes introduce *"Space-based access control with dt.space"*: *"The dt.space attribute is now a primary, permission-relevant field in Grail. Data access can be scoped by space membership, enabling fine-grained governance and consistent data separation across your Grail datasets."* That is a **data-access** control: it sits beside `dt.security_context` in deciding which *records* a team may read.

The gap this column was first drawn for is still open in what Dynatrace documents. Nothing yet delegates which *configuration objects* a team may own — changing an alert, an SLO or a maintenance window still goes through whoever holds the settings permissions, which is how a central team becomes a queue. Whether spaces will also delegate ownership of configuration objects is **not documented yet**: as of 09/28/2026 there is no docs page for spaces. `dt.space` is listed in the published semantic dictionary as an `experimental` permission field (*"The space associated with the record. Spaces enable granular scoping and delegation of monitoring responsibilities."*), but it was not yet in this validation tenant's

`dt.semantic_dictionary.fields` on 09/28/2026 — check yours. That field description's "delegation of monitoring responsibilities" is not a documented mechanism for owning configuration objects. Verify that 1.348 has reached your tenant before relying on `dt.space` ; until then — and for configuration ownership in any case — design against what the other columns give you today.

Sources: [What's new in Dynatrace SaaS 1.348 \(DT docs\)](#) — the `dt.space` release note quoted above (pre-release, read 09/24/2026). [Permission fields — semantic dictionary \(DT docs\)](#) — "*The space associated with the record. Spaces enable granular scoping and delegation of monitoring responsibilities.*" (listed `experimental`). **Dictionary:** no row for `dt.space` in the validation tenant under `filter contains(name, "space")` , read 09/28/2026 (control: the same filter returned 21 other fields).

Inside the OneAgent surface, several distinct concepts share "tag"-adjacent vocabulary. Disambiguating them is essential:

Concept	Field shape	Example	Set how	Notes
Primary fields (<i>sprint-1.337+</i>)	Reserved Dynatrace key (<code>dt.*</code>)	<code>dt.security_context</code> , <code>dt.cost.costcenter</code> , <code>dt.cost.product</code>	OneAgent install / <code>oneagentctl --set-host-tag="dt.security_context=<value>"</code> (the form documented on the June-2026 tags hub; the <code>oneagentctl</code> reference retains the older <code>--set-host-property</code> form); or via OpenPipeline enrichment for non-OneAgent sources	Top-level attributes on every signal at ingest (<i>reserved keys — must be set explicitly; not auto-populated</i>). The recommended Gen3-first surface for boundary, cost, and ownership. Ride into Smartscape, IAM, and OpenPipeline routing without parse processors.

Concept	Field shape	Example	Set how	Notes
Primary tags <i>(sprint-1.337+)</i>	Customer-defined namespace (<code>primary_tags.<key></code>)	<pre>primary_tags.team , primary_tags.environment , primary_tags.app</pre>	OneAgent install / <code>oneagentctl --set-host-tag="primary_tags.<key>=<value>"</code> — the <code>primary_tags.</code> prefix must be written explicitly (it is never added automatically); per-process via the <code>DT_TAGS</code> environment variable	Top-level on every signal. As of June 2026 the namespace is first-class: a dedicated primary-tags docs hub documents <code>primary_tags.<key></code> end to end, including the limit of up to 20 primary tags per host or process — excess tags are silently dropped without a warning. Use for dimensions that don't fit a reserved <code>dt.*</code> key.
Custom host tags <i>(legacy at-source tag)</i>	Plain key/value on the host	<pre>Environment:prod , Owner:platform-team</pre>	OneAgent install / <code>oneagentctl --set-host-tag=<value></code> (no namespace prefix)	Pre-1.337 surface. Still works but lacks the primary-fields propagation guarantees; for new work prefer primary fields/tags.
Auto-tagging rules <i>(legacy, view-time)</i>	Tag conditions on entities (Settings → Tags → Automatically applied tags)	Rule: "if <code>host.name</code> matches <code>^prod-</code> then tag <code>Environment=prod</code> "	Settings 2.0 schema or legacy Configuration API	Computed when an entity is viewed; not present on ingest. Avoid for new work — see §7 for why.

`oneagentctl` syntax — **primary fields vs tags (updated June 2026)**: The dedicated [OneAgent tag setup \(DT docs\)](#) page documents a **single form for both**: `--set-host-tag` carries reserved primary fields and customer primary tags alike — e.g., `oneagentctl --set-host-tag="dt.security_context=confidential"` and `oneagentctl --set-host-`

```
tag="primary_tags.environment=production" ; per-process values use
DT_TAGS="primary_tags.team=bravo" . The older oneagentctl reference still documents --set-host-
property=dt.security_context=easytrade_sec — both pages were live as of 06/10/2026. For new work,
lead with the --set-host-tag form: it is the convention the latest-Dynatrace tags hub builds on.
```

Why primary fields/tags are the recommended Gen3-first default

In recent SaaS sprints, `dt.security_context` and customer-defined `primary_tags.*` became first-class top-level attributes on **every** signal — metrics, spans, logs, business events, Smartscape entities — when set on the OneAgent at install time. AWS Lambda primary-tag propagation followed shortly after via a Lambda-extension environment variable surface (verify the exact variable name and propagation guarantees against current Lambda-extension docs). The `dt.cost.costcenter` and `dt.cost.product` keys are documented as reserved primary Grail fields alongside `dt.security_context` — **reserved keys that require explicit configuration** (`oneagentctl --set-host-tag` or an OpenPipeline enrichment rule); they are not auto-populated. Once set, they are enriched onto every signal at ingest with no further per-signal configuration. `dt.host_group.id` is the exception — it genuinely is auto-enriched from the host-group assignment with no configuration at all.

Three reasons primary fields/tags are the correct default for new work:

1. **They flow without enrichment.** The value lands as a top-level field at ingest. No OpenPipeline parse processor, no auto-tagging rule, no view-time computation needed. DQL filters work directly: `filter dt.security_context == "team-a"`.
2. **They feed all downstream surfaces uniformly.** IAM `MATCH(dt.security_context)` boundary clauses, OpenPipeline `route` rules, Smartscape Ownership, and dashboard filters all see the same value.
3. **They are stable through time.** Set once at OneAgent install, the value rides through every host restart, every process restart, and every signal — no drift between dashboards and alerts.

Why auto-tagging rules are *not* the default

Auto-tagging rules compute a tag at view/query time from entity properties (host name, process name, etc.). They are convenient but introduce three problems:

- **Computed at view time, not ingest** — historical signals that predate a rule change get re-tagged retroactively from the *current* rule, not the rule that was in effect when the signal was generated. This breaks point-in-time analysis.
- **Don't propagate to logs, business events, or spans uniformly** — auto-tags primarily live on entities (hosts, services, process groups). DQL on `fetch logs` won't see the auto-tag without joining through the entity.
- **Couple a tag's value to a regex over a property that may change** — e.g., a rule that derives `Environment` from a host-name prefix locks the team into never renaming hosts.

Per Gen3-first guidance: tag at source via primary fields/tags, not at view time via auto-tagging rules. The migration path from a legacy tenant with auto-tagging rules in place is covered in §6 and §7.

In community practice the sprint-1.337 / sprint-1.338 dates and the AWS Lambda propagation surface above are the commonly-cited values — verify both against your own tenant's release notes before planning around them.

Sources: [Primary Grail fields and tags \(DT docs\)](#) — the propagation and notification-routing statements quoted above; [Configure security context \(DT docs\)](#) — "At-source enrichment is always preferred over OpenPipeline-based enrichment for the security context.", the diagram's layer order; [Primary tags \(DT docs\)](#) — `primary_tags.*` naming convention; primary Grail fields — `dt.host_group.id` (auto-enriched) plus the reserved, explicitly-configured `dt.security_context` , `dt.cost.costcenter` , `dt.cost.product` ; [OneAgent tag setup \(DT docs\)](#) — `--set-host-tag` for both primary fields and tags; "up to 20 primary tags per host or process; excess tags are silently dropped without a warning"; [oneagentctl \(DT docs\)](#) — retains the older `--set-host-property=dt.security_context=easytrade_sec` form.

3. The Four-Source Hierarchy

Breaking (OneAgent 1.345): host `primary_tags` are promoted only from host tags. Verbatim: "Starting with OneAgent version 1.345, the OS Agent derives `primary_tags` from host tags (`hostautotag.conf`) rather than from host properties (`hostcustomproperties.conf`) at the host level." (Re-read 09/02/2026: the release note now reads *derives ... rather than*, where earlier revisions of this entry quoted *promotes ... and no longer from*. The behaviour is unchanged — only Dynatrace's wording moved.) This is a silent failure mode, which is what makes it worth auditing rather than noting: a host that reaches 1.345 simply stops promoting any `primary_tags.*` configured through the properties file, and nothing errors — the dimension just disappears from that host's signals, and every segment, bucket rule, and dashboard keyed on it quietly loses those records. Audit `hostcustomproperties.conf` across the fleet **before** it upgrades and migrate those entries to `hostautotag.conf` (or `oneagentctl --set-host-tag=`). Note the narrow scope: this is about *promotion to primary tags at the host level*, and the `oneagentctl` reference still documents `--set-host-property=` for `dt.*` reserved keys. OneAgent 1.345 released 08/12/2026 with a **staged rollout from 08/25/2026** — tenant version is not agent version, so check the fleet, not the tenant.

Source: [What's new in OneAgent 1.345 \(DT docs\)](#)

Dynatrace's `tags-and-metadata` documentation enumerates seven categories of tags and metadata (auto-detected attributes, manually defined tags, automatically generated tags, cloud-imported tags, Kubernetes / OpenShift labels, Cloud Foundry organization names, and environment-variable-based tags). For practical strategy, this FAQ collapses them to **four operational source buckets** that map cleanly onto where the tag is *set* and how it *propagates* — each bucket appropriate for some dimensions and inappropriate for others.

Source	Where set	Field surface in DQL	Best for
OneAgent — primary fields/tags	Host install / <code>oneagentctl --set-host-tag=</code> for both <code>dt.*</code> reserved keys and explicit <code>primary_tags.<key></code> values (June-2026 tags-hub form; the <code>oneagentctl</code> reference retains <code>--set-host-property=</code> for <code>dt.*</code> keys)	<code>dt.security_context</code> , <code>dt.cost.costcenter</code> , <code>dt.cost.product</code> , <code>primary_tags.<key></code>	Stable, host-level dimensions: security boundary, cost center, environment, team, app — anything that should ride on every signal at source

Source	Where set	Field surface in DQL	Best for
Kubernetes labels and annotations	K8s manifests / Helm / GitOps	<code>k8s.<resource>.label.<key></code> (e.g., <code>k8s.pod.label.app</code> , <code>k8s.namespace.label.team</code>); annotations available as custom metadata at the process level	Dimensions that vary at the pod / workload / namespace level (more granular than the host) — application name, version, environment within a shared cluster
Cloud-provider tags	AWS / Azure / GCP console / IaC tooling	<code>aws.tag.<key></code> , <code>azure.tag.<key></code> , <code>gcp.label.<key></code> plus provider attributes	The source of record for cost-allocation and compliance data when the cloud provider is the canonical owner of that information; useful as input to OpenPipeline enrichment that normalizes them into <code>dt.*</code> primary fields
Auto-tagging rules (legacy)	Settings 2.0 schema	View-time tag conditions on entities	Avoid for new work; acceptable only as a stop-gap on legacy tenants pending migration to primary fields

June 2026 — the tags hub formalizes which sources can emit primary tags directly. The new [Tags documentation hub \(DT docs\)](#) documents primary Grail tags as settable from **five sources**: OneAgent (host tags / `DT_TAGS`), Kubernetes (namespace- or pod-level `metadata.dynatrace.com/primary_tags.<key>` annotations — pod-level values win over namespace-level, and Operator 1.10.0 adds cluster-scoped DynaKube resource attributes below annotations in the specificity chain), cloud-provider tags (AWS / Azure / GCP), OpenTelemetry resource attributes (`OTEL_RESOURCE_ATTRIBUTES`), and OpenPipeline (derived from any incoming field at ingest). Central-configuration rules that promote existing Kubernetes labels to primary tags (limited to 20 rules per scope) ship with **SaaS 1.343** as *Centralized telemetry metadata enrichment* — key-value pairs, namespace annotations, and domain tags managed centrally. SaaS 1.343's rollout started 07/14/2026 with a staged tenant rollout from mid-July 2026 — verify the feature has reached your tenant before relying on it; until it arrives, the at-source mechanisms described here remain the working path. The four-bucket model above still holds — what changes is that each bucket increasingly emits `primary_tags.*` natively instead of relying on enrichment workarounds.

July 2026 — the version floors differ by capability, and your fleet is a separate question again. This entry previously quoted the Kubernetes tag-setup page as requiring "OneAgent version 1.343+, ActiveGate version 1.341+, Dynatrace Operator version 1.10+" for both the at-source and central options. **Corrected 07/31/2026:** that no longer matches the source, and it bundled two capabilities that have different floors.

Capability	Minimum version	Source
At-source / static tag enrichment (host tags, DT_TAGS)	OneAgent 1.333+	OneAgent tag setup (DT docs) , OneAgent attribute enrichment (DT docs)
Central enrichment configuration	OneAgent 1.343	What's new in OneAgent 1.343 (DT docs)
Kubernetes telemetry enrichment	Operator 1.10+, OneAgent 1.333+, ActiveGate 1.343+	Kubernetes tag setup (DT docs)

The practical difference is large: **at-source tagging is available from 1.333, not 1.343**, so a fleet in the 1.333–1.342 band can already emit primary tags at source and is only missing the *central configuration* surface. The ActiveGate figure was also **1.343+**, not 1.341+. Operator 1.10.0 and **OneAgent 1.343 (released 07/28/2026)** have both shipped. What that does *not* mean is that the requirement is satisfied in your environment. **Tenant version is not agent version**, and OneAgent fleets upgrade on their own schedule (auto-update rings, maintenance windows, or manual rollouts), so **verify the installed OneAgent version on the hosts concerned** rather than inferring it from the tenant. Treat any host on **1.342 or earlier as still on the prior mechanisms**, and expect a mixed fleet for weeks — a cluster where some nodes satisfy the requirement and others do not will produce partial primary-tag coverage that looks like a configuration error rather than a version skew.

OneAgent 1.343 is also where the **host/process metadata source** starts to arrive, rather than a single switch flipping. Two 1.343 changes carry it: **Smartscape identifiers are now included in process metadata files**, and **OneAgent can ingest enrichment configuration containing conditional rules expressed as DQL matchers** — which is what lets a rule decide at source whether a given process should carry a value. Verify the agent version per host before designing around either.

Until fleets reach 1.343, the working mechanisms are unchanged: Kubernetes

`metadata.dynatrace.com/primary_tags.<key>` annotations for pod- and namespace-scoped dimensions, and `DT_TAGS / oneagentctl --set-host-tag` (or the older `--set-host-property` form for `dt.*` keys) for host- and process-level dimensions. Neither is superseded by 1.343 — both remain the documented at-source path and continue to work afterwards.

Propagation Depth — Why the Source Matters

Not all sources propagate to all signal types. This table is the load-bearing reason to choose source carefully:

Source	Metrics	Spans	Logs	Business Events	Smartscape Entities
OneAgent primary fields/tags (<i>sprint-1.337+</i>)	✓ at ingest	✓ at ingest	✓ at ingest	✓ at ingest	✓ as primary attribute
K8s labels/annotations	✓ (via DynaKube enrichment)	✓ (via DynaKube + OneAgent)	✓ (via OpenPipeline enrichment)	(rare)	✓ on K8s entity types

Source	Metrics	Spans	Logs	Business Events	Smartscape Entities
Cloud-provider tags	✓ (via cloud integration)	(limited; spans may not carry cloud tags directly)	(limited; depends on log ingestion path)	(limited)	✓ on cloud-resource entity types
Auto-tagging rules	(computed at view)	(computed at view)	(rare; not on raw log records)	(rare)	✓ at view time

If you tag at the OneAgent layer with primary fields, the value is on every signal type at ingest — DQL filters on `dt.security_context` work uniformly across `fetch logs`, `fetch spans`, `fetch bizevents`, `timeseries` queries, and entity queries. If you tag only via cloud-provider tags, the same query returns inconsistent results because the tag is on the cloud entity but not on every signal.

Propagation rule of thumb

For dimensions that need to filter / scope / route across signal types, tag at the OneAgent layer with primary fields/tags. For dimensions that are inherently scoped to a layer (K8s pod, AWS Lambda function), tag at that layer and rely on enrichment to surface them where needed.

Sources:

- ~ [Tags and metadata \(DT docs\)](#) — seven documented categories of tags and metadata; basis for the four operational buckets in this FAQ
- ~ [Tags documentation hub \(DT docs\)](#)
- ~ [Primary tags \(DT docs\)](#) — five documented primary-tag sources
- ~ [Kubernetes tag setup \(DT docs\)](#) — `metadata.dynatrace.com/primary_tags.<key>` annotations, pod-over-namespace precedence, and the minimum component versions, re-read at source 08/24/2026 as **Dynatrace Operator 1.10+**, **OneAgent 1.333+**, **ActiveGate 1.343+** (the "OneAgent 1.343+ / ActiveGate 1.341+" pairing quoted in earlier revisions of this entry is superseded — see the corrected floors above)
- ~ [OneAgent 1.343 release notes \(DT docs\)](#) — released 07/28/2026; Smartscape identifiers in process metadata files, and enrichment configuration with conditional rules expressed as DQL matchers
- ~ [oneagentctl \(DT docs\)](#) — `--set-host-property` vs `--set-host-tag` syntax (classic reference)
- ~ [OpenPipeline \(DT docs\)](#) — enrichment processors that surface K8s and cloud tags as `dt.*` primary fields
- ~ **Derived:** the "verify the agent version per host, treat ≤1.342 as prior-mechanism, expect a mixed fleet" guidance combines the documented version minimums with the fact that agent fleets upgrade independently of tenant version

3.1 Release-identity environment variables (`DT_RELEASE_*`)

One environment-variable family deserves calling out separately, because it is set like a tag, surfaces like a tag, and yet is easy to look for in the wrong place. Release identity is declared on the process via:

Variable	Holds	Example
<code>DT_RELEASE_PRODUCT</code>	Product / component name	<code>easytrade</code>

Variable	Holds	Example
DT_RELEASE_VERSION	Release version	1.5.2
DT_RELEASE_STAGE	Deployment stage	production
DT_RELEASE_BUILD_VERSION	Build identifier, when distinct from the version	4471

Read them from the `tags` map on a Smartscape process node, with *unquoted* bracket keys:

```
smartscapeNodes "PROCESS"
| fieldsAdd product = tags[DT_RELEASE_PRODUCT], version = tags[DT_RELEASE_VERSION]
| filter isNotNull(product) and isNotNull(version)
| dedup {product, version}
```

`tags["DT_RELEASE_PRODUCT"]` — the quoted form — is a **parse error**, not an empty result.

Three traps, all confirmed against a live tenant on 07/30/2026:

- **Do not expect a `deployment.release_*` field on the process node.** The semantic dictionary does define `deployment.release_product`, `deployment.release_version`, `deployment.release_stage`, `deployment.release_build_version`, and `cicd.deployment.release_stage` — but `dt.smartscape.process` declares 30 fields and **none of them is release-related**; the values arrive inside `tags`. Filtering a process node on `deployment.release_version` matches nothing and returns **zero rows rather than an error**, which reads as "no releases tagged" when it actually means "wrong access path."
- **DT_RELEASE_STAGE is frequently unset even where product and version are.** On the validation tenant, 16 processes carried `easytrade / 1.5.2` with `stage` null throughout. Treat stage as optional: filter on product and version, and let stage be a display column rather than a required grouping key.
- **Keep the values slug-safe** — lowercase, hyphenated, no spaces, quotes, `#`, or `&`. This is not cosmetic: these values get concatenated into URLs for deep links into issue trackers, where an unescaped separator silently truncates the target parameter. See **DASH-06 § 7** for that recipe.

3.2 Two tagging domains — and which one wins when both set the same key

The documentation now splits the tagging surface into two named domains, each with its own page. The split matters operationally, because it is the difference between a change you can make in the UI and a change that touches every host.

	Central enrichment	OneAgent domain
Where configured	Rules in Dynatrace	On the host / in the process
Mechanisms	Kubernetes metadata (namespace / workload / pod labels and annotations); cloud-provider tags (AWS, Google Cloud, Azure); host and process properties from OneAgent-monitored infrastructure; custom literal values applied conditionally	Installer flag <code>--set-host-tag</code> ; <code>oneagentctl post-install</code> ; <code>DT_TAGS</code> per process
Deployment cost	<i>"No changes on the hosts are required. No agent restart is needed."</i>	Host access, and a restart for the installer path

	Central enrichment	OneAgent domain
Writes to	<code>dt.cost.product</code> , <code>dt.security_context</code> , <code>dt.cost.costcenter</code> , <code>custom primary_tags.*</code>	The same fields

Custom literal values are the one to notice — a fixed value applied conditionally, centrally, with no host involvement. That covers the common case (stamp `dt.cost.costcenter` on everything matching a condition) without the fleet work the OneAgent domain implies.

Timing. Central rules are not post-processing: *"Enrichment rules run before data enters the processing pipeline, so the same fields are available for pipeline routing, bucket assignment, access control, and cost allocation without any post-processing steps."* That is what makes them interchangeable with at-source tags for OpenPipeline routing and bucket assignment.

Precedence, in two layers. Within central enrichment: *"When rules at multiple scopes target the same dimension for the same signal, the narrower-scope rule takes precedence", and "Rules are evaluated in priority order. When two rules target the same dimension for the same signal, the higher-priority rule wins."*

Across the two domains — the answer to *"I set it in both places, which one lands?"*:

"When the same key is set at multiple scopes, the more specific definition wins:"

1. **Process** — `DT_TAGS`
2. **Host** — `installer` or `oneagentctl`
3. **Ingest enrichment configuration rule**

Read that ordering before debugging a value that "won't change": a central rule cannot override a host tag carrying the same key, so a stale `oneagentctl` value on one host will quietly beat a correct central rule, on that host only. That is a partial-coverage symptom that looks like a broken rule.

Storage shape. *"Enrichment keys are stored under `primary_tags.<key>` , except for `dt.cost.costcenter` , `dt.cost.product` , and `dt.security_context` , which are applied as-is."* So a central rule writing `team` produces `primary_tags.team` , and the three reserved cost/security keys stay top-level — the same shape § 2 describes for the at-source path.

New (SaaS 1.346 — staged rollout from 08/25/2026): classic configuration surfaces can now match on trace-based `primary_tags` . Verbatim: *"Request Naming, Service Detection Rules, Failure Detection and Request Attributes now support trace-based `primary_tags.*` values in the Process group tag condition field."*

This narrows a long-standing split. The `primary_tags.*` values described above could be queried in DQL but could not be used as a condition in those four classic configuration surfaces, which forced a parallel set of process-group tags maintained purely to drive rules. From 1.346 the same enrichment keys can do both jobs — so a `primary_tags.team` established once via the precedence chain above can drive service detection and request naming directly, rather than being mirrored into a second tagging scheme. Worth knowing before you build that mirror; if you already have one, this is the change that lets you retire it.

Sources: [Central enrichment rules \(DT docs\)](#) — the four enrichment sources, the target fields, the before-the-pipeline timing, and both precedence rules, all quoted above, [OneAgent tag setup \(DT docs\)](#) — `--set-host-tag / oneagentctl / DT_TAGS` , the cross-domain precedence order, and the `primary_tags.<key>` storage rule. Read at source 08/27/2026;

What's new in Dynatrace SaaS 1.346 (DT docs) — trace-based `primary_tags.*` in the Process group tag condition field, quoted above.

4. AWS / Azure / GCP — Per-Cloud Specifics

Cloud-provider tags are the source of record for cost-allocation, compliance, and ownership data when the cloud provider owns the resource lifecycle. They land in Dynatrace via the **Clouds app** — modern, with **AWS new-connections GA**, **Azure new-connections GA**, and **GCP new-connections** documented as “*will follow soon*” (classic GCP integration is GA) — or via legacy integrations (CloudWatch monitor for AWS, Azure Monitor for Azure, GCP integration for GCP).

AWS

AWS surface	Field in Dynatrace	Notes
EC2 / RDS / Lambda / ECS / EKS resource tags	<code>aws.tag.<TagKey></code>	Tag keys preserved verbatim, including casing — <code>aws.tag.CostCenter</code> $\neq$ <code>aws.tag.costcenter</code>
Account ID	<code>aws.account.id</code>	Always present on AWS-sourced signals
Account alias	<code>aws.account.alias</code>	When set in IAM
Region	<code>aws.region</code>	Programmatic region code (e.g., <code>us-east-1</code>)
Resource ARN	<code>aws.arn</code>	Available on resources with discoverable ARNs

AWS Lambda primary-tag propagation: In community practice, recent Dynatrace Lambda-extension releases surface a `DT_*` environment-variable mechanism that propagates AWS-resource attributes (such as `aws.arn`, `aws.region`, `aws.account.id`) and customer-defined primary tags into the Lambda's logs and spans, not just metrics — closing the gap where serverless functions previously had partial tag propagation. Verify the exact variable name and propagation surface against the [current Dynatrace OpenTelemetry-on-Lambda documentation](#), since the Lambda-extension surface area evolves sprint-to-sprint (the docs prominently surface `DT_OPEN_TELEMETRY_ENABLE_INTEGRATION` rather than a single `DT_TAGS` envelope).

AWS integration path: prefer the **Clouds app** for new tenants — direct cloud connection without ActiveGate, GA for AWS. The legacy CloudWatch monitor remains supported but does not benefit from continued enhancement.

Azure

Azure surface	Field in Dynatrace	Notes
Resource tags (VMs, App Services, AKS, etc.)	<code>azure.tag.<TagKey></code>	Tag keys preserved verbatim
Resource group	<code>azure.resource_group</code>	Useful as a coarse boundary if not tagged explicitly

Azure surface	Field in Dynatrace	Notes
Subscription	<code>azure.subscription_id</code> , <code>azure.subscription_name</code>	Always present on Azure-sourced signals
Region	<code>azure.region</code>	Programmatic naming (in a recent SaaS sprint). Region strings render as lowercase with no spaces (e.g., <code>eastus2</code> , not <code>East US 2</code>). Update DQL filters and dashboards comparing <code>azure.region</code> against display strings — use <code>in(azure.region, {"eastus2", "westus2"})</code> . Verify the exact rollout sprint against your tenant's release notes.

Azure integration path: Clouds app (preview) is the modern path; Azure Monitor integration via ActiveGate remains the production-stable option.

GCP

GCP surface	Field in Dynatrace	Notes
Resource labels (Compute Engine, GKE, Cloud Run)	<code>gcp.label.<LabelKey></code>	GCP enforces lowercase + hyphens on label keys; values still arbitrary
Project	<code>gcp.project_id</code>	Always present on GCP-sourced signals
Region / zone	<code>gcp.region</code> , <code>gcp.zone</code>	Programmatic codes

GCP integration path: Clouds app (preview); legacy GCP integration via ActiveGate also supported.

Cross-Cloud Tag Naming Drift — A Concrete Problem

The same conceptual dimension lands with different keys across providers:

Concept	AWS	Azure	GCP
Cost center	<code>aws.tag.CostCenter</code>	<code>azure.tag.costCenter</code>	<code>gcp.label.cost_center</code>
Environment	<code>aws.tag.Environment</code>	<code>azure.tag.environment</code>	<code>gcp.label.environment</code>
Team / owner	<code>aws.tag.Owner</code>	<code>azure.tag.owner</code>	<code>gcp.label.team</code>
Application	<code>aws.tag.Application</code>	<code>azure.tag.app</code>	<code>gcp.label.application</code>

The casing alone (`CostCenter` vs `costCenter` vs `cost_center`) means a naive query has to enumerate every variant. Most teams discover this when they try to build a single "cost by team" dashboard that has to span clouds — and the dashboard is full of `coalesce(...)` that papers over the inconsistency.

The recommended fix is to normalize at ingest, not at query time. OpenPipeline enrichment processors can map `aws.tag.CostCenter` / `azure.tag.costCenter` / `gcp.label.cost_center` into a single canonical `dt.cost.costcenter` field, so DQL queries and IAM policies see one consistent surface regardless of cloud provenance. See OPIPE topic series for the worked enrichment-processor examples.

Sources:

- [Clouds app \(DT docs\)](#)
- [OpenTelemetry on AWS Lambda \(DT docs\)](#) — Lambda extension and `DT_*` integration env-var surface
- [OpenPipeline \(DT docs\)](#)
- [Tagging AWS resources \(AWS general reference\)](#) — "*Tag values are case sensitive*" and "*tag keys are case sensitive*"; basis for the cross-cloud casing-drift problem
- [Tag resources \(Azure docs\)](#) — "*Tag names are case insensitive for operations*", while "*Tag values are case sensitive*"; 50-tag limit
- [Best practices for resource labels \(Google Cloud docs\)](#) — GCP label naming rules (lowercase, hyphens) and recommended dimensions

5. Tagging Standards — Taxonomy and Naming

A tag *strategy* is the source-of-truth and propagation decisions in §6. A tag *standard* is the taxonomy and naming convention that makes the strategy executable. Standards before tools — pick the dimensions and the names before configuring anything.

Recommended Dimensions to Tag

Most tenants benefit from tagging on these seven dimensions. Not every dimension needs to be a primary field; some can stay as ordinary tags or labels. The point is to decide *which* dimensions matter and to name them consistently.

Dimension	Purpose	Recommended source	Recommended key
Environment	Separate prod / nonprod / dev for thresholds, alerts, IAM scope	OneAgent primary tag at install	<code>primary_tags.environment</code> (values: <code>prod</code> , <code>nonprod</code> , <code>dev</code>)
Application / service	Identify which app a host or pod serves	OneAgent primary tag (host-level) or K8s label (pod-level)	<code>primary_tags.app</code> (host) or <code>k8s.pod.label.app</code> (pod)
Team / owner	Who owns this; drives notification routing	OneAgent primary tag at install	<code>primary_tags.team</code>
Cost center	FinOps attribution; cost-allocation reports	OneAgent primary field (preferred) or normalized from cloud tag via OpenPipeline	<code>dt.cost.costcenter</code>
Product / business line	Higher-level grouping above cost center	OneAgent primary field	<code>dt.cost.product</code>

Dimension	Purpose	Recommended source	Recommended key
Security context	IAM boundary for record-level / field-level access (the canonical Gen3 boundary)	OneAgent primary field at install	<code>dt.security_context</code>
Compliance / criticality	Regulatory or business-criticality scoping	OneAgent primary tag	<code>primary_tags.compliance</code> (values: <code>pci</code> , <code>pii</code> , <code>sox</code> , <code>none</code>) or <code>primary_tags.criticality</code> (<code>tier1</code> , <code>tier2</code> , <code>tier3</code>)

Optional further dimensions: deployment region, data residency, lifecycle phase. Add them only when there is a concrete consumer (a dashboard, an alert routing rule, an IAM policy) that would use them.

Reserved Dynatrace Keys

Some keys are reserved by Dynatrace. Don't redefine these — use them as-is, or use a different key:

Reserved namespace	Owned by	Don't use for
<code>dt.*</code>	Dynatrace platform (e.g., <code>dt.security_context</code> , <code>dt.cost.costcenter</code> , <code>dt.cost.product</code> , <code>dt.host.*</code> , <code>dt.smartscape.*</code>)	Custom dimensions — use <code>primary_tags.<key></code> instead
<code>aws.*</code> , <code>azure.*</code> , <code>gcp.*</code>	Cloud-provider integrations	Custom dimensions; these are populated by the integration
<code>k8s.*</code>	Kubernetes integration	Custom dimensions; these are populated by DynaKube
<code>host.*</code> , <code>process.*</code> , <code>service.*</code>	OneAgent semantic dictionary	Custom dimensions

Naming Conventions

Element	Convention	Example	Anti-example
Key casing	lowercase, kebab-case if multi-word	<code>primary_tags.cost-center</code> (or use <code>dt.cost.costcenter</code>)	<code>primary_tags.CostCenter</code> , <code>primary_tags.cost_center</code> (<i>pick one and stick to it</i>)
Value casing	lowercase, kebab-case	<code>prod</code> , <code>nonprod</code> , <code>team-payments</code>	<code>Prod</code> , <code>NonProd</code> , <code>Team-Payments</code>
Value stability	use values that don't change frequently	<code>team-payments</code> , <code>app-checkout</code>	<code>release-2026-q2</code> , <code>incident-12345</code> (<i>metadata of the day</i>)
Spaces	never	<code>team-payments</code>	<code>team payments</code>

Element	Convention	Example	Anti-example
Reserved values	avoid <code>null</code> , <code>none</code> , <code>default</code> , empty string	<code>none-set</code> if you must	<code>null</code> , <code>""</code>
Cross-cloud normalization	one canonical key per dimension, regardless of provider	All cost-center tags resolve to <code>dt.cost.costcenter</code>	<code>aws.tag.CostCenter</code> and <code>azure.tag.costCenter</code> both queried separately forever

Industry Frameworks Worth Reading

If your organization doesn't have a tagging standard, these are reasonable starting points to adapt:

- **AWS Well-Architected — Tagging Best Practices** ([docs](#)) — practical taxonomy guidance including the 7 dimensions above
- **Azure — Develop your naming and tagging strategy** ([docs](#))
- **GCP — Best practices for resource labels** ([docs](#))
- **FinOps Foundation — Tagging best practices** — for the cost-attribution dimensions specifically

These frameworks broadly agree on the seven dimensions table above. Use them as the starting point and adapt names to your organization's existing conventions where they exist.

Sources:

- ~ [AWS Well-Architected — Tagging Best Practices](#) — taxonomy guidance and the canonical seven-dimension surface
- ~ [Azure cloud-adoption-framework — naming and tagging](#)
- ~ [GCP — Best practices for resource labels](#)
- ~ [Kubernetes — Labels and Selectors](#) — K8s label key constraints (DNS subdomain, optional prefix, length)
- ~ [FinOps Foundation](#)

6. Strategy Themes

Standards say *what* to call things. Strategy says *who owns the value*, *what wins when sources conflict*, and *what to do when a source is missing*. Five themes:

6.1 One Source of Truth Per Dimension

In community practice, the most workable arrangement is to designate exactly one authoritative source per dimension.

Dimension	Authoritative source	Why
Cost center	OneAgent primary field (or AWS Tag Editor for cloud-only resources, normalized via OpenPipeline)	One person's spreadsheet, not three

Dimension	Authoritative source	Why
Environment	OneAgent primary tag at install	Stable per host, must be on every signal type
Team / owner	OneAgent primary tag at install	Tied to host placement, not pod scheduling
App name	OneAgent primary tag (host-level) or K8s label (pod-level) — pick one based on your topology	Mixing the two creates ambiguous joins
Security context	OneAgent primary field at install	Must be tamper-resistant; can't be a view-time computation

Mixing sources for the same dimension is the most common cause of tagging drift. If `team` is sometimes on the host and sometimes on the K8s pod label, every dashboard has to decide which to trust per-row.

6.2 Precedence and Conflict Resolution

When two sources both carry the same dimension, document which wins. A simple, stable precedence:

1. **OneAgent primary field/tag** (set explicitly at install) wins — most stable, most propagated
2. **K8s label** (when scoping inherently to pod/namespace) — next
3. **Cloud-provider tag** (when the cloud is the source of record for that resource) — next
4. **Auto-tagging rule** — last resort, only on legacy tenants

Document the precedence in your runbook. The runbook is consulted when an auditor asks "why does this report show team-A but the IAM policy targets team-B?"

6.3 Cardinality Control

Not every cloud tag should become a Dynatrace primary tag. Cloud accounts often accumulate dozens of tags per resource — automation tags, cost-allocation tags, compliance tags, deployment-pipeline tags. Surfacing them all into Dynatrace creates Smartscape clutter, dashboard-filter sprawl, and IAM-policy churn.

Curate which dimensions propagate as primary. A reasonable default: only the seven dimensions in §5 land as primary fields/tags. Other cloud tags can remain queryable as `aws.tag.<key>` / `azure.tag.<key>` / `gcp.label.<key>` for the ad-hoc queries that need them, without elevating them to first-class status.

6.4 Bare-Metal and On-Prem Fallback

Cloud tags don't exist on bare-metal hosts, on-prem VMs, or air-gapped infrastructure. **OneAgent host tags / primary tags are the universal floor** — they work the same way on every host regardless of provider.

Strategy implication: if your tenant spans cloud + on-prem, the *primary* tagging surface must be the OneAgent layer. Cloud-provider tags become a *secondary* enrichment for cloud-resident workloads. The opposite arrangement (cloud-tags-as-primary) leaves on-prem hosts with no consistent tag value.

6.5 Source-Side Enrichment Over View-Time Rules

When a dimension's authoritative source isn't already a Dynatrace-shaped key (e.g., `aws.tag.CostCenter` rather than `dt.cost.costcenter`), do the normalization **at ingest** via OpenPipeline enrichment processors, not

at view time via auto-tagging rules.

Reasons:

- The normalized value lands on every signal at ingest (rather than being computed per-query)
- The OpenPipeline rule lives in source-controlled config (Terraform / Monaco / GitOps)
- A change to the rule is point-in-time; historical signals carry the value that was correct at the time they were ingested
- IAM policies and Smartscape see the normalized value uniformly

See OPIPE topic series for the worked enrichment patterns. The migration path from a legacy tenant: identify the auto-tagging rules that compute primary-dimension values, replace them with OpenPipeline enrichment, then deprecate the auto-tagging rules.

When the source of truth is an external CMDB (rather than a field already on the incoming signal), OpenPipeline-at-ingest doesn't apply — the CMDB values aren't on the data. The fit there is a scheduled workflow that reconciles the CMDB onto host tags at source: **WFLOW-08 §11 (CMDB-Driven Host Tag Enrichment)** provides an import-ready template that reads CMDB lookup tables and sets `dt.security_context` / `dt.cost.costcenter` / `primary_tags.*` via the OneAgent Remote Configuration Management API, with dry-run guardrails.

Sources: [oneagentctl \(DT docs\)](#) — primary-field assignment surface that anchors the precedence model in §6.2, [OpenPipeline \(DT docs\)](#) — source-side enrichment as the alternative to view-time auto-tagging rules (§6.5), [Host groups \(DT docs\)](#) — host-group-level tag rideability as the universal floor (§6.4).

7. Anti-Patterns

Patterns that work in the short term and create rework in the long term. Each is followed by the recommended alternative.

7.1 Auto-Tagging Rule Sprawl

Symptom: the Settings → Tags → Automatically applied tags page has dozens of rules, often with overlapping conditions, that compute primary dimensions (environment, team, cost center) from regex over `host.name` or `process.name`.

Why it's a problem: rules are computed at view time; they don't propagate to logs / business events / spans uniformly; changing a rule retroactively re-tags historical signals; rules couple a tag's value to a property (host name) that may need to change for unrelated reasons.

Alternative: tag at source via OneAgent primary fields/tags. For cloud workloads, normalize cloud-provider tags via OpenPipeline enrichment. Migrate auto-tagging rules incrementally — replace one rule at a time, verify the at-source value matches the computed value, then disable the rule.

7.2 Treating Every Cloud Tag as Primary

Symptom: the Clouds app integration is configured to surface all `aws.tag.*` / `azure.tag.*` / `gcp.label.*` keys as Dynatrace primary tags. Smartscape is cluttered. Dashboard filter dropdowns show 80+ tag keys.

Why it's a problem: cloud accounts accumulate operational tags (deployment pipelines, automation jobs, ticket numbers) that aren't observability-meaningful. Surfacing them all elevates noise to first-class status.

Alternative: curate the dimensions that propagate as primary (the seven in §5). Other cloud tags remain queryable on `aws.tag.<key>` / `azure.tag.<key>` / `gcp.label.<key>` without first-class status.

7.3 Inconsistent Casing Across Clouds Without Normalization

Symptom: dashboards and DQL queries paper over `aws.tag.CostCenter` vs `azure.tag.costCenter` vs `gcp.label.cost_center` with `coalesce(...)`. The same dimension is queried differently in every report.

Why it's a problem: every new dashboard re-derives the normalization logic. Inconsistencies multiply. New team members don't know which key to query.

Alternative: normalize at ingest with OpenPipeline enrichment processors. One canonical `dt.cost.costcenter` field, regardless of cloud provenance.

7.4 Cost Center via Host-Name Regex

Symptom: an auto-tagging rule derives `CostCenter` from a host-name pattern like `^(?<cc>[a-z]{4})-prod-`.

Why it's a problem: hosts get renamed. Naming conventions evolve. The rule breaks silently. Worse — if the rule is wrong, the wrong cost center gets reported, often for months before someone notices the column doesn't sum to the total cloud bill.

Alternative: set `dt.cost.costcenter` explicitly at OneAgent install, sourced from the same authoritative system (CMDB, FinOps spreadsheet, AWS Tag Editor) as the cost-allocation report.

7.5 Relying on Cloud Tag for IAM Boundary Without Verifying Propagation

Symptom: an IAM policy uses `MATCH(aws.tag.SecurityContext)` as the boundary, assuming the AWS tag flows to every signal type the policy needs to scope.

Why it's a problem: AWS tags don't uniformly propagate to logs, business events, or all entity types. The policy works for some queries and silently fails-open for others. The user sees more data than the policy intended.

Alternative: use `MATCH(dt.security_context)` set at the OneAgent layer (or normalized via OpenPipeline from the cloud tag). The boundary then propagates to every signal type uniformly. See IAM topic series for the boundary-standardization pattern.

7.6 Tagging Metadata of the Day Into Long-Lived Primary Fields

Symptom: primary fields carry values like `release-2026-q2`, `incident-12345`, `feature-flag-payment-v2`. New values get added every sprint.

Why it's a problem: primary fields are designed for stable values that ride on every signal forever. Short-lived values churn the value-set, break dashboard filter dropdowns, and pollute IAM policy match conditions.

Alternative: primary fields/tags carry stable dimensions only (env, team, app, cost center, security context). Short-lived metadata (release version, incident ID, feature flag) belongs on **business events** — `fetch bizevents | filter event.release == "2026-q2"` works without polluting the host's primary tag set.

Sources: [Tags and metadata \(DT docs\)](#) — auto-tagging-rule semantics and view-time computation behavior, [Host groups \(DT docs\)](#) — naming-constraint anti-patterns (cannot start with `dt.`, 100-character maximum).

8. Final Recommendation

Tagging is foundational, not cosmetic. A coherent strategy pays back across every Dynatrace surface — Smartscape, dashboards, alerts, IAM, OpenPipeline, automation, cost attribution.

Four principles, in order of priority:

1. **Tag at source.** OneAgent primary fields/tags at install time are the universal floor. Cloud-provider tags supplement; auto-tagging rules retire.
2. **Use Dynatrace primary fields/tags as the canonical surface.** `dt.security_context`, `dt.cost.costcenter`, `dt.cost.product`, `primary_tags.<key>` — these are what dashboards, IAM, and routing should query against.
3. **Normalize cloud-provider tags via OpenPipeline enrichment.** Cross-cloud dimensions (cost center, environment, team) land in canonical `dt.*` fields at ingest, not at query time.
4. **Standards before tools.** Decide the seven dimensions, the naming convention, and the source-of-truth per dimension *before* configuring OneAgent installers, OpenPipeline processors, or IAM policies.

The cost of a bad tagging strategy is paid at every dashboard, every audit, every incident, every cost-allocation report — for as long as the tenant exists. The cost of a good one is paid once, up front.

Summary

Tagging in Dynatrace draws from four sources (OneAgent, Kubernetes, cloud-provider integrations, and legacy auto-tagging rules) that propagate differently and serve different purposes. Primary fields and primary tags (sprint-1.337+) are the recommended Gen3-first surface for stable dimensions because they ride on every signal at ingest. Standards (taxonomy + naming) come before strategy (source-of-truth + precedence + cardinality + fallback + source-side enrichment). The canonical pattern: tag at source via OneAgent, normalize cloud tags via OpenPipeline enrichment, and use `dt.*` primary fields as the surface every consumer queries.

Next Steps

- Inventory the dimensions your tenant currently tags on (across all four sources)
- Identify mismatches: dimensions tagged in multiple sources, dimensions tagged inconsistently across clouds, primary dimensions computed via auto-tagging rules
- Define the seven-dimension standard (or your organization's adapted version) before the next host onboarding wave
- Plan the OpenPipeline enrichment processors that normalize cross-cloud tags into `dt.*` canonical fields
- Cross-reference with FAQ-01 (host group naming strategy) and the IAM / ORGNZ / OPIPE topic series for the consuming patterns

Sources: the §8 Final Recommendation is a **Derived** synthesis from §1–§7 — no single source endorses the combined "primary fields/tags first, K8s/cloud tags as enrichment, OpenPipeline normalization at ingest, auto-tagging only as legacy stop-gap" recommendation as a single statement. The Sources blocks on §1–§7 list the inputs the synthesis rests on.

Related Resources

This FAQ does not stand alone — tagging strategy is decided alongside host-group naming, IAM boundary design, OpenPipeline routing, and per-cloud integration mechanics. Per-section citations live in the `<sub>Sources: </sub>` blocks above; this section points at the wider reading list.

Companion FAQ entries:

- **FAQ-01: Why you need a good Host Group naming strategy** — host-group boundaries; tagging-source-of-truth and host-group naming are decided together
- **FAQ-03: OneAgent vs OpenTelemetry — A Decision Framework** — primary-field tagging assumes OneAgent presence; FAQ-03 covers when OneAgent is and is not the right tool

Topic series in this collection (the consuming patterns):

- **OPIPE series** — OpenPipeline enrichment processors for cross-cloud tag normalization and primary-field assignment at ingest
- **CLOUD series** — per-provider integration deep dives (AWS, Azure, GCP) and field mapping reference
- **IAM series** — `MATCH(dt.security_context)` boundary patterns; the canonical Gen3 ABAC surface for record-level access
- **AUTOM series** — config-as-code for any remaining tag rules (Terraform / Monaco / GitOps)
- **K8S series** — Kubernetes label propagation through DynaKube and OneAgent metadata enrichment
- **ORGNZ series** — segments and bucket strategy that consume tags as scoping inputs

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official [Dynatrace documentation](#).