

FAQ-08: How Does OneAgent Decide Which Logs to Collect?

Series: FAQ — Frequently Asked Questions | **Reference:** 08 — How OneAgent Decides Which Logs to Collect | **Created:** June 2026 | **Last Updated:** 09/28/2026

Overview

A common surprise after deploying OneAgent is that *some* log files show up in Dynatrace automatically and others don't — with no obvious pattern from the outside. The application log under `/opt/app/logs/` appears within a minute; the one your batch job writes to `/data/exports/run.out` never does. Neither was configured by hand, so why the difference?

The answer is **log auto-discovery** — a mechanism inside OneAgent that continuously scans the host, decides which files are logs worth collecting, and starts ingesting them. It is not magic and it is not arbitrary: it applies an explicit, documented set of rules. Understanding those rules turns "why isn't my log here?" from a guessing game into a three-question checklist, and tells you exactly when to reach for a **custom log source** instead.

This entry explains how the decision is made, what gets picked up automatically versus what needs a nudge, and how to verify what was actually discovered in your environment.

Table of Contents

1. [Short Answer](#)
2. [The Mental Model — Three Gates](#)
3. [The Auto-Discovery Requirements](#)
4. [The Built-in Include / Exclude Rules](#)
5. [What's Auto-Detected vs. What Needs Help](#)

6. [When Auto-Discovery Misses Your Log — Custom Log Sources](#)
7. [Scale and Limits](#)
8. [Recommended Approach](#)
9. [Common Gotchas](#)

Prerequisites

Requirement	Details
Audience	Platform/infra engineers, SREs, and onboarding leads deciding what log coverage to expect from OneAgent and where manual configuration is needed
Format	Decision-support document — explains the discovery mechanism and the auto-vs-custom decision; no hands-on lab
Deployment	Dynatrace SaaS with Grail; Log Monitoring powered by the OneAgent log module (enabled by default). Applies to host OneAgent on Linux and Windows
Related topic series	OPLOGS (OpenPipeline log processing), OPMIG (Classic Log Monitoring → OpenPipeline), ONBRD-07 (validating discovered data), K8S (container/Kubernetes log collection), S2D (locating logs during migration)
Related FAQ	FAQ-03 (OneAgent vs OpenTelemetry — includes Windows Event Log and IIS handling)

1. Short Answer

OneAgent runs a **log auto-discovery engine** that re-evaluates the host roughly **every 60 seconds**. A file is collected automatically only when it passes **three independent gates**:

Gate	The question it answers	Pass condition
Process	<i>Who is holding this file?</i>	The file is kept open, in write mode, by an important (deep-monitored) process


Gate	The question it answers	Pass condition
Pattern	<i>Where is it and what is it named?</i>	Its path/name matches a built-in include rule and no exclude rule
File properties	<i>What kind of file is it?</i>	Text content, supported encoding, recent activity, above a minimum size, append-only

If a file fails any gate, auto-discovery skips it — and the fix is almost always a **custom log source** (§6), not a tweak to the built-in rules (which you cannot expand, only narrow).

By default OneAgent auto-discovers logs from a wide range of technologies: **process log files, system logs (syslog), journald, IIS logs, Windows Event Log channels, and container stdout / stderr**. Log Monitoring is on by default; your **log ingest rules** then decide which of the discovered records are actually stored in Grail.

Sources: [Log autodiscovery \(DT docs\)](#) — discovery runs "every 60 seconds"; [Log ingestion via OneAgent \(DT docs\)](#) — default-detected source list and ingest-rule gating. **Derived:** the "three gates" framing is a synthesis of the requirements, built-in rules, and file constraints documented separately on the autodiscovery page — Dynatrace documents the conditions but does not group them this way.

2. The Mental Model — Three Gates



Reading the gates in order is the fastest way to diagnose a missing log.

Gate 1 — Process. Auto-discovery is process-anchored, not filesystem-anchored. OneAgent finds candidate files by looking at what its **deep-monitored processes** have open, not by walking the whole disk. The docs put it plainly: *"The log file must be kept open by an important process."* "Important" here means a process Dynatrace deep-monitors — one belonging to a known technology or consuming significant resources. This single fact explains most surprises: a log written by a short-lived cron job, a process Dynatrace doesn't deep-monitor, or a file nothing holds open will not be discovered, no matter where it lives.

Gate 2 — Pattern. Among the files that survive Gate 1, OneAgent applies a fixed set of **built-in include/exclude rules** (§4) against the path and filename. Files in a `log / logs` directory, or whose names carry a `log` token, are included; database transaction logs, container runtime directories, and Windows event-trace files are explicitly excluded.

Gate 3 — File properties. Finally, the file itself must look like an active text log: supported encoding, recently written, above a minimum size, appended-to rather than rewritten. The full list is in §3.

A file is collected only when it clears **all three**. The gates are independent — passing two of three is not enough.

Sources: [Log autodiscovery \(DT docs\)](#) — "The log file must be kept open by an important process"; [How OneAgent works \(DT docs\)](#) — definition of important / deep-monitored processes. **Derived:** the ordering of the three gates as a diagnostic sequence is an authoring synthesis, not a documented procedure.

3. The Auto-Discovery Requirements

For a file to be auto-discovered, **all** of the following must hold. These are the literal conditions from the autodiscovery reference:

#	Requirement	Detail
1	Held open by an important process	The file must be kept open by a process OneAgent deep-monitors
2	Minimum age	The file must have existed for at least one minute

#	Requirement	Detail
3	Supported encoding	UTF-8 by default; also UTF-8 BOM and, if the byte-order mark is present, UTF-16LE and UTF-16BE
4	Recent activity	The file must have been written to within the last 7 days
5	Path or name match	The file must be in a <code>log / logs</code> folder (or a subfolder), or its filename must contain a <code>log</code> string preceded or followed by a period (<code>.</code>) or underscore (<code>_</code>)
6	Recent timestamps	Records are ingested only if their timestamp is within the last 24 hours ; entries timestamped more than 10 minutes ahead of the current time are overridden
7	Security rules	The file path must satisfy OneAgent's security rules

And the file must behave like a live, append-only text log:

- It **cannot be deleted earlier than a minute** after creation.
- It must be **appended** to (old content is not rewritten).
- It must have **text content** (binaries are not auto-detected — see §5).
- It must be **opened constantly**, not just briefly while a line is written.
- It must be **opened in write mode**.
- It must be **at least the configured size threshold** (default **500 bytes**).

Rolling out (OneAgent 1.343) — the rapid-rotation constraint narrows on Linux.

OneAgent 1.343 (released 07/28/2026) adds support for **rapidly rotated and compressed log files on Linux**, which is precisely the case the "cannot be deleted earlier than a minute after creation" bullet and the one-minute minimum age (#2) exclude today. Two limits on that relief matter more than the headline:

- **It is Linux-only.** On Windows the constraint above is unchanged, and a rapidly rotated file there still needs a custom log source (§6).
- **It depends on the agent version, not the tenant version.** Fleets upgrade on their own schedule, so **verify the OneAgent version on the hosts concerned**; anything on 1.342 or earlier behaves exactly as the table describes. Expect a mixed fleet for weeks.

Until the hosts you care about are on 1.343, treat requirement #2 and the one-minute-deletion bullet as fully in force — they remain the working model.

Two of these trip people up most often. The **7-day activity window** (#4) and the **24-hour timestamp rule** (#6) are *different* checks: #4 is about when the file was last *written*, #6 is about the *timestamps inside the records*. A file actively written today but full of records dated last week will be discovered yet have its old records dropped by the timestamp rule. And the **path/name rule** (#5) is why `/data/exports/run.out` from the Overview is invisible — it is neither under a `log / logs` path nor does its name carry a `log` token.

Sources: [Log autodiscovery \(DT docs\)](#) — requirements list, encodings, 7-day activity window, 24-hour timestamp rule, 10-minute future-timestamp override, append-only constraints, and the 500-byte default size threshold (all quoted from the page), [OneAgent 1.343 release notes \(DT docs\)](#) — released 07/28/2026; support for rapidly rotated and compressed log files on Linux. **Derived:** the "verify the agent version per host, expect a mixed fleet" qualifier follows from agent fleets upgrading independently of tenant version.

4. The Built-in Include / Exclude Rules

Gate 2 is driven by a fixed set of **built-in autodetector rules**. Each rule matches an operating system, a directory pattern, and a file pattern, and either **includes** or **excludes** the match. Exclude rules win where they overlap with include rules — which is how database and container-runtime paths stay out even though they may sit near `log` directories.

OS	Directory pattern	File pattern	Action
All	<code>/Log/</code>	<code>*.xel</code>	Exclude
All	<code>/commitlog/</code>	<code>*.log</code>	Exclude
Windows	<code>/CCM/Logs/</code>	<code>*</code>	Exclude
Windows	<code>/MSSQL/Log/</code>	<code>*.trc</code>	Exclude
Windows	<code>/*MSSQL*/OLAP/Log/</code>	<code>*.trc</code>	Exclude
Windows	<code>MSSQL/DATA/</code>	<code>*.ldf</code>	Exclude
Windows	—	<code>*.evtx</code>	Exclude
Linux	<code>/var/log/pods/</code>	<code>*</code>	Exclude

OS	Directory pattern	File pattern	Action
Linux	/var/lib/docker/containers/	*	Exclude
All	/	*[-. _]log	Include
All	/	catalina.out*	Include
All	/log/	*	Include
All	/log/*/	*	Include
All	/logs/	*	Include
All	/logs/*/	*	Include
Linux	^/var/log/**/	*	Include

What this tells you:

- ***[-. _]log is the filename escape hatch.** A file does not have to live under a `log / logs` directory — `payments-2026.log` , `audit_log` , or `app.log` anywhere on disk matches the include rule by name. (`catalina.out*` is a special-cased Tomcat include.)
- **Database transaction logs are deliberately excluded** (`*.xel` , `/commitlog/` , `*.trc` , `*.ldf`). These are not application logs and would be noise; that exclusion is by design, not a gap.
- ***.evtx is excluded as a file** because Windows Event Log is collected through the Event Log API, not by reading the `.evtx` files (see §5).
- **/var/log/pods/ and /var/lib/docker/containers/ are excluded** because Kubernetes and Docker container logs are collected via the dedicated container-log path, not as host files — excluding them here prevents double-collection.

Crucially, the security rules can only **narrow** this set — *"the custom security rules can only narrow auto-detection but not expand it."* There is no setting that adds new include rules. To collect a file the built-in rules don't match, you use a custom log source (§6).

Sources: [Log autodiscovery \(DT docs\)](#) — built-in autodetector rules table and the "security rules can only narrow auto-detection but not expand it" constraint (both quoted from the page). **Derived:** the "why each exclusion exists" annotations combine the rule table with the container-collection and Windows-Event-Log handling documented elsewhere on the page.

5. What's Auto-Detected vs. What Needs Help

"Which logs" is broader than just files held open by a process. By default OneAgent auto-collects several distinct source types:

Source	Auto-detected?	Notes
Process / application log files	Yes	The core file-discovery path described in §2–§4
System logs (syslog)	Yes	Enabled by default
journald	Yes	systemd journal on Linux
IIS logs	Yes (Windows)	On by default; can be turned off in advanced OneAgent log settings if not needed
Windows Event Log	Yes (Windows)	System / Application / Security channels via the Event Log API — <i>not</i> by reading <code>.evtx</code> files (which the rules exclude)
Container stdout / stderr	Yes (Linux)	In Kubernetes, OpenShift, and non-instrumented Docker, via the container-log path — not the host file rules

And several things are deliberately **not** auto-collected, each with a specific reason:

- **Binary log files** — *"Binary log files are not detected automatically."* You must add a custom log source with the binary-format option enabled.
- **Logs on network filesystems (Linux)** — detection of NFS-mounted logs is **disabled by default** and must be explicitly enabled.
- **OneAgent's own logs** — collected only if you enable *"Allow OneAgent to monitor Dynatrace logs."*
- **Files that fail any gate in §3** — wrong path/name, no holding process, stale, too small, non-text, or rewritten-in-place rather than appended.

The log-module configuration lives under **Settings** → **Collect and capture** → **Log monitoring** → **Configure log module**; the network-filesystem option and *Allow OneAgent to monitor Dynatrace logs* are on its **Advanced settings** page. Container and Kubernetes specifics (the container-log path, namespace scoping, masking) are covered in depth in the K8S series rather than here.

Sources:

- [Log ingestion via OneAgent \(DT docs\)](#) — default-detected sources (system logs, IIS, Windows Event Log, containers)
- [Log autodiscovery \(DT docs\)](#) — binaries not auto-detected, network-filesystem default-off, "Allow OneAgent to monitor Dynatrace logs"; "Settings > Collect and capture > Log monitoring > Configure log module > Advanced settings"
- [Advanced log settings \(DT docs\)](#) — per-technology toggles (e.g., turning off IIS log detection)
- [Windows event logs \(DT docs\)](#) — System / Application / Security channel collection

6. When Auto-Discovery Misses Your Log — Custom Log Sources

When a file fails a gate, the fix is a **custom log source** — an explicit path you tell OneAgent to monitor. A custom log source **supplements** auto-detection for a specific path; it does **not** change the built-in rules or expand what auto-discovery finds elsewhere.

Where to configure (narrower scope wins) — the page is the same at every scope; for host and host group, first use **Go to scope** in the upper-left corner of Settings to select the host or host group:

Scope	Path
Environment	Settings → Collect and capture → Log monitoring → Configure log module → Sources
Host group	Settings → Go to scope (host group) → Collect and capture → Log monitoring → Configure log module → Sources
Host	Settings → Go to scope (host) → Collect and capture → Log monitoring → Configure log module → Sources

In each case, select **New log source rule** in the *Add missing log sources* section.

Path rules:

- Paths must be **absolute** — Windows `letter:\...`, Linux `/...`, NFS `//hostname/...`. Relative paths are rejected.
- Up to **100 paths per rule**, and up to **1000 rules per scope**.

- Wildcards: `*` substitutes any string of characters *except* `/` or `\` (valid in both directories and filenames); `#` substitutes a string of digits (valid in **filenames only**).
Example: `/var/log/app/access_#.log` Or `/opt/*/out/*.log`.
- The OneAgent account (`dtuser` by default) needs **read** permission on the file; on remote drives it also needs read+execute on each directory.
- Each custom path must still satisfy OneAgent's **security rules**.

When you specifically need a custom source:

- The file lives outside a `log / logs` path and its name has no `log` token (the Overview's `/data/exports/run.out`).
- Nothing keeps the file open, or the writing process isn't deep-monitored.
- The file is **binary** (enable the binary-format option).
- The file uses an unusual **rotation pattern** the detector doesn't follow cleanly — **narrowed on Linux from OneAgent 1.343** (released 07/28/2026), which adds support for rapidly rotated files. On **Windows**, or on any host still running **1.342 or earlier**, this remains a live cause and the custom-source fix still applies. Check the agent version on the host before concluding the rotation pattern is the problem: tenant version is not agent version.
- The file is **compressed** (`.gz` and similar). Compressed logs are not part of the ordinary auto-discovery text-file model — **from OneAgent 1.343 they are collectable on Linux**; before that version, and on Windows, a compressed file needs either a custom source with the appropriate handling or a change to the writer so an uncompressed live file exists to tail. Again, verify the agent version deployed.

Because custom sources don't expand the rules, the durable fix for "we have a whole class of logs in a non-standard location" is sometimes simpler: write those logs to a `log / logs` directory, or give them a `*.log` name, so the built-in include rules catch them without per-host configuration.

Sources: [Custom log source \(DT docs\)](#) — config scopes and precedence ("*Go to Settings > Collect and capture > Log monitoring > Configure log module > Sources*"), absolute-path requirement, `*` / `#` wildcard semantics, 100 paths/rule and 1000 rules/scope, `dtuser` permission requirement, "custom log sources do not expand auto-detection", [Log autodiscovery \(DT docs\)](#) — binary-format option for binary files; rotation-pattern handling, [OneAgent 1.343 release notes \(DT docs\)](#) — released 07/28/2026; support for rapidly rotated and compressed log files on Linux. **Derived:** the "write to a log/

directory or use a .log name" recommendation combines the §4 include rules with the custom-source-doesn't-expand constraint — neither source states it as advice.

7. Scale and Limits

Auto-discovery is bounded so it can't overwhelm a busy host:

Limit	Value
Files per directory the log module handles	up to 10,000
New log content per minute	up to 200 MB
Default maximum log sources per process-group instance	200
Minimum file size to be discovered	500 bytes (default threshold)
Paths per custom log source rule	100
Custom log source rules per scope	1,000

These matter at design time. A directory with tens of thousands of rotated files, or a process group that legitimately emits more than 200 distinct log sources, will hit a ceiling — and the symptom (some files silently not collected) looks identical to a discovery rule miss. If you operate at that scale, split logs across directories, raise the relevant maximums where the settings allow, and verify coverage explicitly (§8) rather than assuming everything is captured.

Sources: [Log autodiscovery \(DT docs\)](#) — 10,000 files per directory, 200 MB/min, default 200 log sources per process-group instance, 500-byte threshold; [Custom log source \(DT docs\)](#) — 100 paths/rule and 1,000 rules/scope.

8. Recommended Approach

Don't assume coverage — confirm it, then patch the gaps deliberately.

1. **Verify what was actually discovered before adding anything — from the log-module self-monitoring events.** The OneAgent log module emits a self-monitoring

event per discovered source into `dt.system.events` under the `Log Module` provider. Each `log_source.status` event carries the source's current file status and ingest status, so one events query gives you the whole coverage picture — including sources that were **detected but are not being ingested**, which a `fetch logs` query cannot show (no stored records means nothing to count):

```
fetch dt.system.events, from:-24h
| filter event.provider == "Log Module" and event.type == "log_source.status"
| dedup event.id, sort:{ timestamp desc }
| filter event.status == "Active"
| fieldsAdd coverage =
  if(log.source.file_status == "FILE_STATUS_OK" and log.source.ingest_status
== "Ingested", "Available | Ingesting", else:
    if(log.source.file_status == "FILE_STATUS_OK" and log.source.ingest_status
== "Not ingested", "Available | NOT ingesting", else:
      if(log.source.ingest_status == "Partially ingested", "Partial ingestion",
        else: "Unsupported or issue")))
| summarize sources = count(), by:{coverage}
| sort sources desc
```

`log.source.file_status` reflects the source's current state (`FILE_STATUS_OK` , `FILE_STATUS_IS_BINARY` , `FILE_STATUS_NOT_EXIST` , ...) and `log.source.ingest_status` is one of `Ingested` / `Not ingested` / `Partially ingested` . Crossing the two gives a coverage matrix — the cell that matters most is **available but not ingesting**: a source OneAgent can read but no ingest rule keeps.

	Ingesting	NOT ingesting
Available (<code>FILE_STATUS_OK</code>)	✅ working as intended	⚠️ detected, silently uncollected — the gap to investigate
Unsupported / issue	ingesting despite a file issue	❌ dead — neither cleanly readable nor collected

To list the sources behind the ⚠️ cell, swap the `summarize` for a filter and `fields` :

```

fetch dt.system.events, from:-24h
| filter event.provider == "Log Module" and event.type == "log_source.status"
| dedup event.id, sort:{ timestamp desc }
| filter event.status == "Active"
| filter log.source.file_status == "FILE_STATUS_OK" and
log.source.ingest_status == "Not ingested"
| fields log.source, log.source.file_status, log.source.ingest_status,
log.source.origin
| limit 50

```

This surface is billing-cheap — it scans events, not raw log records (0 bytes billed in testing) — and needs only event-read permission rather than logs-table read. It is delivered in **Early Access** and requires opt-in: the OneAgent log module must be enabled to send self-monitoring events to your tenant, and its content is folding into the built-in **Log ingest Overview** dashboard for Dynatrace **1.339+** (staged rollout — verify the events are flowing in your tenant before relying on them). Until they are, the Grail log-records query below remains the working path.

Fallback (no opt-in / older tenant). Query the stored log records directly. This shows only sources that produced records, so it cannot separate "detected but not ingested" from "never detected" — but it needs no opt-in. Group by `log.source` alone: `log.source.file_status` and `log.source.ingest_status` are fields of the `log_source.status` self-monitoring events above, not of log records, and are empty on every log record (0 of 20,006,675 records over 24 h on the validation tenant, 09/28/2026):

```

fetch logs, from:-24h
| summarize records = count(), by:{log.source}
| sort records desc
| limit 20

```

The **Discovery & Coverage** view in the platform gives the same picture interactively.

2. **For a missing log, walk the three gates in order (§2).** Is a deep-monitored process holding it open? Does its path/name match an include rule and dodge the exclude rules? Does the file itself meet the §3 properties? The first failing gate is your answer.
3. **Prefer making the file discoverable over per-host config.** If you control the application, writing to a `log / logs` directory or using a `*.log` name lets the built-in rules collect it everywhere with zero configuration — more durable than a custom source maintained per host.
4. **Use custom log sources for what genuinely can't be moved** — binaries, fixed non-standard paths, files no monitored process holds open. Set them at the broadest scope

that's correct (environment > host group > host) to minimize drift.

5. **Don't try to "open up" auto-discovery.** Security rules only narrow it; there is no expand knob. Custom sources are the supported way to add coverage.
6. **Mind the storage step.** Discovery is not ingestion. Log Monitoring is on by default, but your **log ingest rules** decide what's stored in Grail and what's dropped — a source can be discovered and still produce no stored records because a rule excludes it.

Sources: [Log autodiscovery \(DT docs\)](#) — narrow-only security rules; [Log ingestion via OneAgent \(DT docs\)](#) — ingest rules gate storage. The `dt.system.events / Log Module / log_source.status` coverage queries and the `log.source.file_status` (`FILE_STATUS_OK` / `FILE_STATUS_IS_BINARY` / `FILE_STATUS_NOT_EXIST`) and `log.source.ingest_status` (`Ingested` / `Not ingested` / `Partially ingested`) values were executed live against a tenant (07/08/2026) — the "available but not ingesting" cell returned real sources (OneAgent's own logs, `/var/log/syslog`). The self-monitoring event stream is Early Access, opt-in, and folds into the built-in *Log ingest Overview* dashboard at Dynatrace 1.339+.

Derived: the file-status × ingest-status coverage matrix and the verify-first / make-discoverable-before-custom ordering are authoring recommendations, not a documented procedure.

9. Common Gotchas

Symptom	Likely cause	Where
Log file exists but never appears	Path isn't under <code>log / logs</code> and name has no <code>log</code> token	§3 #5, §4
File appears but old lines are missing	Records' timestamps are outside the 24-hour window	§3 #6
Nothing from a short-lived job's log	No long-lived deep-monitored process keeps the file open	§2 Gate 1
Database <code>.trc</code> / <code>.ldf</code> / <code>commitlog</code> not collected	Excluded by built-in rules by design	§4
<code>.evtx</code> files "ignored" on Windows	Event Log is read via the API, not the files	§5

Symptom	Likely cause	Where
Container logs missing under <code>/var/log/pods</code>	Collected via the container path, not host file rules	§4, §5
A binary log is invisible	Binaries aren't auto-detected	§5, §6
Logs on an NFS mount not collected (Linux)	Network-filesystem detection is off by default	§5
Source shows as discovered but no records in Grail	A log ingest rule is dropping it (discovery $\neq$ storage)	§8
Some files in a huge directory collected, others not	Hit the 10,000-files / 200 MB-per-minute / 200-sources ceiling	§7

Sources: [Log autodiscovery \(DT docs\)](#), [Log ingestion via OneAgent \(DT docs\)](#), [Custom log source \(DT docs\)](#) — each row maps to the requirement, rule, source-type, or limit cited in the referenced section.

 **Disclaimer:** This content is AI-generated, community-driven, and **not supported by Dynatrace**. Log auto-discovery behavior, built-in rules, and limits evolve across OneAgent versions — always verify against the current [Dynatrace documentation](#) for your environment before relying on a specific rule or threshold.