

MCH-06: Capacity and Scaling

Series: MCH — Managed Cluster Health | **Notebook:** 6 of 8 | **Created:** September 2026 | **Last Updated:** 09/25/2026

Overview

Status: Outline — this notebook's content is in development. The sections below list what it will cover.

Layer 3 of the health model — whether the cluster can keep up, including with a node gone. Read MCH-01 first for the cluster architecture and the five-layer health model this notebook builds on.

Table of Contents

1. [Node Sizing](#)
 2. [Headroom for Failure](#)
 3. [Capacity Signals](#)
 4. [Adding Nodes](#)
 5. [Summary and Next Steps](#)
-

Prerequisites

Requirement	Details
Dynatrace deployment	A self-hosted Dynatrace Managed cluster. Nothing in this series applies to Dynatrace SaaS.
Access	Cluster Management Console (CMC) administrator access; for the Cluster API, a cluster API token with the <code>ServiceProviderAPI</code> permission
Query language	None. Managed has no Grail, so this series has no DQL cells — every check is a CMC view, a cluster event, or a Cluster API call
Version basis	Written against the Managed documentation read 09/25/2026, when 1.342 – 1.346 were the generally supported releases

1. Node Sizing

Planned content:

- The Micro–XLarge sizing table: vCPUs, RAM, host units, storage per store

2. Headroom for Failure

Planned content:

- The one-third-above-typical rule and what it means for 3, 5 and 7 nodes

3. Capacity Signals

Planned content:

- *Adaptive Load Reduction activity*, the *Cluster health self-monitoring* dashboard, `dsfm:` metrics

4. Adding Nodes

Planned content:

- Scale-out rules: same hardware, ≤ 10 ms inter-node latency, 30-node maximum

5. Summary and Next Steps

Content in development. For the health model and the full series lineup, see MCH-01.

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official [Dynatrace documentation](#).