

S2S-99: Best Practice Summary

Series: S2S — SaaS to SaaS Migration | **Notebook:** 99 | **Created:** March 2026 | **Last Updated:** 09/28/2026

Overview

This notebook consolidates **112 actionable best practices** from the S2S series into a single reference organized around the 9-step migration framework. Each practice is definitive: it tells you exactly what to set, when, and why. Use this as a pre-flight checklist and ongoing reference throughout your SaaS-to-SaaS migration.

The S2S series follows a **9-step framework** (Discover through Optimize) supported by an **11-step Order of Operations** that governs execution sequence.

The 9-Step Migration Framework

Step	Name	Focus
1	Discover	Migration scenarios and inventory
2	Strategize	Define your migration approach
3	Design	Target tenant architecture
4	Prepare	Export and pre-stage
5	Execute	Configuration import and agent cutover
6	Integrate	Cloud, dashboards, and workflows
7	Expand	OpenPipeline, SLOs, and alerting
8	Enable	Parallel operation and stakeholder handover
9	Optimize	Cutover validation and decommission

The 11-Step Order of Operations

#	Operation	Description
1	Assess	Inventory source tenant
2	Provision	Target SaaS tenant and access (SSO/IAM)
3	Export	Configuration from source tenant
4	Migrate	Configuration to target tenant
5	Rebuild	Dashboards, SLOs, alerts
6	Redirect	OneAgents and operators to target
7	Reconnect	Cloud integrations and extensions
8	Migrate	Any remaining configuration
9	Validate	Data flow and performance
10	Cutover	Full switch to target tenant
11	Decommission	Source tenant

OneAgent Attribute Enrichment (1.331+): OneAgent can enrich all telemetry (metrics, spans, logs, events) with primary fields (`dt.security_context` , `dt.cost.costcenter`) and primary tags (`primary_tags.environment` , `primary_tags.team`) at the source. More efficient than auto-tags — feeds directly into OpenPipeline routing, bucket assignment, and Grail permissions. Configure via `oneagentctl --set-host-tag` or `--set-host-tag` at install time. See [docs](#).

Table of Contents

1. [Step 1: Discover — Migration Scenarios and Inventory](#)
2. [Step 2: Strategize — Define Your Migration Approach](#)
3. [Step 3: Design — Target Tenant Architecture](#)
4. [Step 4: Prepare — Export and Pre-Stage](#)
5. [Step 5: Execute — Configuration Import and Agent Cutover](#)
6. [Step 6: Integrate — Cloud, Dashboards, and Workflows](#)
7. [Step 7: Expand — OpenPipeline, SLOs, and Alerting](#)

8. [Step 8: Enable — Parallel Operation and Stakeholder Handover](#)
9. [Step 9: Optimize — Cutover Validation and Decommission](#)
10. [The Five Principles of S2S Migration](#)

Prerequisites

Requirement	Details
S2S Series	Read S2S-01 through S2S-09 for full context
Source Tenant	Active Dynatrace SaaS with admin access
Target Tenant	Provisioned Dynatrace SaaS with admin access
CLI Tools	Monaco CLI v2.x, Terraform v1.5+ with provider ~> 1.91

1. Step 1: Discover — Migration Scenarios and Inventory

Source: S2S-01

#	Best Practice	Recommended Setting/Value	Priority
1	Complete entity discovery before anything else	Run <code>fetch dt.entity.host \ summarize count()</code> for hosts, services, process groups, and applications to establish baseline numbers	Critical
2	Identify your migration scenario	Classify as consolidation (many-to-one), split (one-to-many), regional relocation, or cloud provider change — each has distinct tooling and planning requirements	Critical
3	Select Monaco for configuration, Terraform for IAM	Monaco v2 covers all 8 config types (settings, document, automation, bucket, segment, slo-v2, openpipeline, classic api); Terraform is required exclusively for IAM policies, groups, and bindings	Critical
4	Plan for entity ID changes	Entity IDs (HOST-xxx, SERVICE-xxx) are tenant-specific and will change; identify every dashboard, SLO, and alert that hardcodes an entity ID	Critical
5	Document the 90/10 manual items	90% of config migrates automatically; budget 90% of your effort for the remaining 10% — entity ID remapping, integration repointing, IAM redesign, credential recreation	Critical

#	Best Practice	Recommended Setting/Value	Priority
6	Build a comprehensive endpoint inventory	Document every agent URL, API call, webhook, and integration endpoint that references the source tenant URL	Critical
7	Triage detected problems before migration	Active problems (especially frequent/duplicate events) carry noise to the target tenant; suppress or tune anomaly detection for sources with >500 active problems	Critical
8	Identify configuration debt to leave behind	Catalog stale maintenance windows, disabled notification rules, inactive synthetic monitors, unused management zones — do not migrate these	Critical
9	Inventory Extensions 2.0 separately	Monaco does not export extension installations; Terraform can install/activate (<code>dynatrace_hub_extension_active_version</code>) and configure (<code>dynatrace_hub_extension_v2_config</code>) Hub extensions — otherwise reinstall from Hub; list all extensions and plan their path before migration	Recommended
10	Audit configuration changes from last 30 days	Query audit logs to identify recently changed settings that may not be in your last export	Recommended

2. Step 2: Strategize — Define Your Migration Approach

Source: S2S-02

#	Best Practice	Recommended Setting/Value	Priority
11	Choose phased approach for environments with >500 hosts	Migrate dev first (weeks 1-2), staging (weeks 3-4), production (weeks 5-6); gives Dynatrace Intelligence time to build baselines	Critical
12	Define success criteria before starting	Document measurable targets: host count parity, service count parity, log volume parity, SLO evaluation accuracy, alerting coverage	Critical
13	Engage your account team for licensing early	Contact Dynatrace to discuss temporary parallel licensing; negotiate before the migration window opens	Critical

#	Best Practice	Recommended Setting/Value	Priority
14	Plan a 2-4 week parallel operation period	Historical data does not migrate; parallel is the only path to data continuity and Dynatrace Intelligence baseline stability	Critical
15	Establish a configuration freeze window	No settings modifications in the source tenant during cutover; communicate freeze dates to all teams	Critical
16	Migrate multi-source consolidations sequentially	Complete Source 1 → validate → Source 2 → validate; never migrate two sources in parallel; migrate the simpler source first (non-K8s before K8s)	Critical
17	Align migration timing with stable traffic periods	Avoid holidays, sales events, or other atypical traffic patterns that would distort Dynatrace Intelligence baselines	Recommended
18	Expect the 90/10 rule	90% of config migrates automatically; budget 90% of effort for the remaining 10%	Critical
19	Prepare a rollback plan	Document how to revert agents to the source tenant if issues arise during cutover	Recommended

3. Step 3: Design — Target Tenant Architecture

Source: S2S-03

#	Best Practice	Recommended Setting/Value	Priority
20	Redesign IAM during consolidation	Use the migration as the opportunity to fix group sprawl and overprivileged policies; redesign from scratch for the target tenant	Critical
21	Deploy data foundation before governance	Buckets and enrichment rules must exist before OpenPipeline routing, segments, SLOs, and IAM policies that reference them	Critical
22	Replace entity IDs with tag-based filters	Convert <code>entityId("HOST-xxxx")</code> to <code>tag(app:checkout), tag(env:production)</code> in dashboards, SLOs, and notification rules before migration	Critical
23	Design bucket naming for consolidation	Use prefix naming (e.g., <code>us-east_app_logs</code> , <code>eu-west_app_logs</code>) to prevent collisions when consolidating multiple tenants	Recommended

#	Best Practice	Recommended Setting/Value	Priority
24	Plan the 24-item deployment order	Buckets (1st) → Enrichment (2nd) → OpenPipeline (3rd) → Segments (4th) → Gen2 config (5th-13th) → Alerting/SLOs (14th-18th) → Dashboards/Workflows (19th-21st) → IAM last (22nd-24th)	Critical
25	Replace management zone ID references with names	Convert <code>mzId(123)</code> to <code>mzName("Production")</code> because zone IDs change between tenants	Critical
26	Prefix resources with source identifier during consolidation	Set <code>name = "\${source_prefix} - \${resource_name}"</code> on all dashboards, rules, and SLOs to prevent naming collisions	Recommended
27	Tag everything by business unit before splitting a tenant	Apply auto-tags identifying ownership before splitting one tenant into many	Recommended
28	Map cloud provider identity constructs when changing providers	AWS IAM roles to Azure Service Principals, Azure AD Groups to GCP Google Groups, etc.	Recommended

4. Step 4: Prepare — Export and Pre-Stage

Source: S2S-04, S2S-10

#	Best Practice	Recommended Setting/Value	Priority
29	Export via <code>monaco download</code>	Run <code>monaco download manifest.yaml --environment source-tenant</code> to export all configuration in one operation	Critical
30	Use the S2S-10 migration scripts for SUA packaging	Automates Monaco download + SUA-compatible <code>.tar.gz</code> packaging; Bash and PowerShell versions provided	Recommended
31	Package exports as <code>.tar.gz</code> for SUA — never <code>.zip</code>	The SaaS Upgrade Assistant rejects <code>.zip</code> archives; must be <code>.tar.gz</code> with <code>exportMetadata.json</code> at root	Critical

#	Best Practice	Recommended Setting/Value	Priority
32	Store exported configuration in Git	Commit the Monaco export directory to a Git repository immediately after download	Critical
33	Validate export completeness	Compare <code>ls projects/full-export/settings/ \</code> <code>wc -l</code> against Settings API <code>totalCount</code> per schema	Critical
34	Validate export contains no secrets	Run <code>grep -r "dt0c01" projects/</code> after export and confirm 0 matches	Critical
35	Run <code>monaco deploy manifest.yaml --environment <target> --dry-run</code> before any deploy	Checks manifest structure, references and template rendering; it does not contact the tenant, so a deploy can still fail with HTTP 400. Monaco has no <code>validate</code> command	Critical
36	Provision target tenant and verify admin access	Confirm SSO, API token creation, and environment admin role before proceeding	Critical
37	Configure SSO with full SAML message signing	Create a new SAML application in your IdP for the target tenant — new Entity ID, new ACS URL; never reuse the source SAML app	Critical
38	Test SSO with a pilot user before cutover	SAML configuration issues are the #1 day-of-cutover blocker	Critical
39	Deploy ActiveGates in the target tenant	ActiveGates cannot be reconfigured like OneAgents; install fresh from the target tenant UI	Critical
40	Prepare K8s operator manifests for the target tenant	Update DynaKube API to <code>v1beta5</code> or <code>v1beta6</code> ; use Operator v1.8.1+ from <code>oci://public.ecr.aws/dynatrace/dynatrace-operator</code>	Critical
41	Recreate network zones in the target	Export with <code>monaco download ... --settings-schema builtin:networkzones</code> and deploy to the target before agent migration	Critical
42	Create new OAuth clients and API tokens in the target	Source credentials cannot be exported; create fresh clients with matching scopes	Critical

#	Best Practice	Recommended Setting/Value	Priority
43	Limit Azure AD group claims to 150 groups per SAML assertion	Azure AD has a hard limit; use group filtering if your org exceeds this	Recommended

5. Step 5: Execute — Configuration Import and Agent Cutover

Source: S2S-05

#	Best Practice	Recommended Setting/Value	Priority
44	Follow the 24-item deployment order	Buckets → Enrichment → OpenPipeline → Segments → Gen2 config → Alerting/SLOs → Dashboards/Workflows → IAM last	Critical
45	Deploy IAM groups, policies, and bindings last	IAM WHERE clauses reference schemas, buckets, and security contexts that must already exist	Critical
46	Use <code>monaco deploy manifest.yaml --environment target</code> for import	Monaco resolves dependencies automatically within a single run (except IAM)	Critical
47	Read a clean dry-run as structure-only	Payload errors (HTTP 400) and missing permissions (HTTP 403) surface only at deploy time — deploy to a non-production target first	Critical
48	Deploy IAM via Terraform with <code>account-idm-read</code> , <code>account-idm-write</code> , and <code>iam-policies-management</code> scopes	These three OAuth scopes are required to manage groups, policies, and bindings	Critical
49	Reconfigure OneAgent with <code>oneagentctl</code> — do not reinstall	Run <code>oneagentctl --set-server=<target-url>/communication --set-tenant=<target-id> --set-tenant-token=<target-token> --restart-service</code> , then restart applications monitored with deep code modules	Critical

#	Best Practice	Recommended Setting/Value	Priority
50	Route firewall-restricted hosts through ActiveGates	Use <code>--set-server="https://<ag-host>:9999/communication"</code> for on-prem or DMZ hosts; semicolon-separate multiple AGs for failover	Critical
51	Restart application processes after agent reconfiguration	Application processes must restart for the agent to report to the new tenant	Critical
52	Migrate non-production first	Dev agents first (weeks 1-2), staging (weeks 3-4), production last (weeks 5-6)	Critical
53	Validate after each migration wave	Verify host count, service count, and data flow in the target after each environment wave	Critical
54	Automate fleet-wide <code>oneagentctl</code> with config management	Use Ansible, Chef, Puppet, AWS SSM, or Azure Automation for large fleets	Recommended
55	Budget 540m CPU and 872Mi memory per node for Dynatrace overhead	OneAgent (100m/512Mi) + CSI Driver (440m/360Mi) per node; control plane adds 650m/320Mi per cluster	Recommended
56	Update IRSA (AWS), Managed Identity (Azure), or Workload Identity (GCP) to reference target	Cloud-specific K8s identity must point to the new tenant	Critical

6. Step 6: Integrate — Cloud, Dashboards, and Workflows

Source: S2S-06

#	Best Practice	Recommended Setting/Value	Priority
57	Recreate cloud provider credentials in the target tenant	<code>monaco download</code> cannot export AWS, Azure, or K8s credentials; create new IAM roles, service principals, and service accounts	Critical

#	Best Practice	Recommended Setting/Value	Priority
58	Create new AWS IAM role with trust policy for the target tenant	Reference the target tenant's Dynatrace account ID and external ID from the AWS integration wizard	Critical
59	Create new Azure App Registration with Monitoring Reader role	New Tenant ID, Client ID, Client Secret for the target tenant	Critical
60	Create new GCP Service Account with Monitoring Viewer and Compute Viewer roles	Generate new JSON key for the target tenant integration	Critical
61	Deploy new AWS Log Forwarder Lambda pointing to the target	Update S3 event notifications to trigger the new Lambda function	Critical
62	Update webhook URLs that reference source-tenant-specific endpoints	Verify custom webhooks accept traffic from the target tenant's IP range	Critical
63	Migrate dashboards via Monaco and audit for hardcoded entity IDs	Search dashboard JSON for <code>entityId(</code> , <code>HOST-</code> , <code>SERVICE-</code> , <code>PROCESS_GROUP-</code> patterns and replace with tag-based filters	Critical

#	Best Practice	Recommended Setting/Value	Priority
64	Reassign dashboard ownership to a target tenant user or service account	Source user accounts may not exist in the target; set <code>owner</code> field to a valid target identity	Critical
65	Create new workflow service user (actor) in the target tenant	Source service users cannot be migrated; create new service user and set as workflow actor	Critical
66	Recreate Synthetic private locations in the target tenant	Private Synthetic locations are tenant-specific; update location references in all monitor configs	Critical
67	Use a classic API token (not OAuth) for Synthetic monitor migration	Synthetic monitors require classic token as of provider v1.88.0+	Critical
68	Plan Extensions 2.0 separately	Monaco does not export extension installations; Terraform can install/activate (<code>dynatrace_hub_extension_active_version</code>) and configure (<code>dynatrace_hub_extension_v2_config</code>) Hub extensions — otherwise reinstall from Hub. Remote extensions run on an ActiveGate group; SQL extensions can alternatively run in Kubernetes via Dynatrace Operator — verify per extension	Critical
69	Standardize tag naming across cloud providers	Use consistent keys (<code>app</code> , <code>env</code> , <code>team</code>) regardless of AWS/Azure/GCP	Recommended
70	Verify every webhook integration (Teams, Slack, PagerDuty, ServiceNow, Jira)	Check each integration's URL and credential; many are unchanged, but none is guaranteed (see #62)	Recommended

7. Step 7: Expand — OpenPipeline, SLOs, and Alerting

Source: S2S-07

#	Best Practice	Recommended Setting/Value	Priority
71	Deploy Grail buckets before OpenPipeline routing rules	Routing rules that target nonexistent buckets cause data loss	Critical
72	Deploy in order: buckets → enrichment rules → OpenPipeline rules	This three-step sequence ensures data flows correctly from the first byte	Critical
73	Wait for K8s enrichment propagation before deploying routing rules	Enrichment rules typically take ~45 min to propagate on EKS; routing rules that depend on enriched attributes fail silently until propagation completes — data lands in <code>default_logs</code>	Critical
74	Add a CI/CD validation gate between enrichment and routing deploys	Query for the enriched field your routing matches on (e.g. <code>isNull(k8s.namespace.name)</code>) before applying routing rules; do not rely on <code>monaco deploy</code> alone (it does not wait for propagation)	Recommended
75	Use prefix naming for bucket consolidation	E.g., <code>us-east_app_logs</code> , <code>eu-west_app_logs</code> to prevent collisions when consolidating tenants	Recommended
76	Match retention policies exactly for regulated environments	Set to same days as source (e.g., 360 days for PCI); verify compliance requirements before cutover	Critical

#	Best Practice	Recommended Setting/Value	Priority
77	Update bucket references in OpenPipeline routing rules	Routing rules from export contain source bucket names; replace with target bucket names	Critical
78	Convert SLO entity IDs to tag-based filters	Replace <code>type(SERVICE), entityId(SERVICE-xxx)</code> with <code>type(SERVICE), tag(app:checkout), tag(env:production)</code>	Critical
79	Set SLO evaluation windows to 30 days initially	Start with <code>ROLLING_3_DAYS</code> , expand to <code>ROLLING_WEEK</code> after 7 days, then to full window after 30 days	Critical
80	Align SLO cutover with calendar boundaries	If using calendar-based SLOs, switch at the start of a month or quarter to avoid partial evaluation	Recommended
81	Port alerting profile severity and tag filters as-is	Ensure the referenced tags exist in the target tenant before deploying alerting profiles	Critical
82	Verify maintenance window timezone settings	Cron expressions are timezone-sensitive; confirm schedules are correct after import	Recommended
83	Update maintenance window entity scopes	Source entity references in scope definitions must be remapped to target tenant entities	Critical

8. Step 8: Enable — Parallel Operation and Stakeholder Handover

Source: S2S-08

#	Best Practice	Recommended Setting/Value	Priority
84	Run parallel operation for 2-4 weeks minimum	Historical data does not migrate; parallel is the only path to continuity	Critical
85	Allow 7-14 days for Dynatrace Intelligence baselines to stabilize	Availability: 2-3 days; response time and error rate: 1-2 weeks; resource utilization: 2-4 weeks (full business cycle)	Critical
86	Use phased parallel model to control cost	Each host reports to one tenant (OneAgent cannot dual-report), so agent cost does not double; overlap cost comes from integrations, synthetics and forwarders you point at both tenants, plus the per-wave window	Critical
87	Communicate to all personas	Engineering teams: 4 weeks before cutover; SRE/on-call: 2 weeks before; executives: weekly status	Critical
88	Communicate the SLO baseline gap to stakeholders	SLOs start at 0% in the target and accumulate data over 7-30 days depending on evaluation window	Critical
89	Warn about Dynatrace Intelligence false positives	Baselines are relearning during first 2-4 weeks; expect noisy alerting until stabilized	Recommended
90	Negotiate dual licensing with Dynatrace	Contact your account team to discuss temporary parallel licensing to reduce cost	Recommended
91	Configure SCIM provisioning for the target tenant	Set new SCIM URL and token; users sync automatically on first login	Recommended
92	Export problem history before decommission	Problem history cannot migrate; export critical incidents as documentation for post-mortem reference	Recommended
93	Keep a fallback local admin account	If SAML is misconfigured on day of cutover, local admin bypasses SSO	Critical

9. Step 9: Optimize — Cutover Validation and Decommission

Source: S2S-09

#	Best Practice	Recommended Setting/Value	Priority
94	Complete go/no-go checklist before cutover	Verify host count, service count, log volume, span volume, SLO evaluation, alerting coverage — all must match success criteria	Critical
95	Validate all data flows via DQL	Run <code>fetch logs, from:-1h \ summarize count() , fetch spans, from:-1h \ summarize count() ,</code> and entity count queries to confirm parity	Critical
96	Verify enrichment tags are populating	Run <code>fetch logs, from:-1h \ filter isNotNull(dt.security_context) \ summarize count() and confirm count > 0</code>	Critical
97	Run final <code>monaco download → monaco deploy sync</code> before cutover	Captures any configuration changes made since initial export	Critical
98	Get stakeholder sign-off	SRE, dev, security, management, and compliance each validate their domain: alerting, dashboards, IAM, SLO reporting, retention	Critical
99	Disable alerting in source tenant after cutover	Prevents duplicate alerts during the transition tail	Critical
100	Keep source tenant running 30 days minimum after cutover	Stop ingest (redirect agents and integrations, revoke ingest tokens); data stays queryable for its bucket retention. Do not deprovision until all stakeholders confirm	Critical
101	Rotate all credentials within 30 days of cutover	API tokens, OAuth clients, cloud provider keys — revoke any temporary migration credentials	Critical
102	Revoke all source tenant API tokens and deactivate OAuth clients	Prevents unauthorized access to the deprecated tenant	Critical
103	Remove SAML/SSO application from IdP for source tenant	Eliminates stale IdP configuration and prevents confusion	Recommended
104	Export final audit logs from source before decommission	Compliance requires audit trail retention; export before access is lost	Critical

#	Best Practice	Recommended Setting/Value	Priority
105	Update all documentation	Replace source tenant URLs, API endpoints, and dashboard links in runbooks, CI/CD pipelines, and architecture diagrams	Critical
106	Tune Dynatrace Intelligence anomaly detection during weeks 1-2 post-cutover	Suppress known false positives; adjust sensitivity for noisy services	Critical
107	Expand SLO evaluation windows to full duration by week 4	Move from <code>ROLLING_3_DAYS</code> to <code>ROLLING_WEEK</code> to <code>ROLLING_MONTH</code> as data accumulates	Recommended
108	Consolidate redundant Synthetic monitors	If consolidating tenants, deduplicate monitors that tested the same endpoints from different tenants	Optional
109	Replace all remaining entity ID references	Post-migration cleanup ensures configuration is fully portable for future migrations	Recommended
110	Schedule credential rotation for cloud provider integrations	Rotate within 30 days of cutover for all AWS, Azure, and GCP integrations	Recommended
111	Conduct a lessons-learned review within 2 weeks of cutover	Capture planning gaps, tool issues, communication effectiveness, timeline accuracy	Recommended
112	Stop ingest into the source tenant after cutover	Redirect agents and integrations and revoke ingest tokens; historical data stays queryable for its bucket retention without active monitoring cost	Critical

The Five Principles of S2S Migration

1. **Export early, validate often** — do not wait for cutover to discover missing config
2. **Tags over entity IDs** — make all configuration portable before migration
3. **Parallel is not optional** — historical data does not migrate; 30 days minimum
4. **IAM last** — deploy the resources first, then lock them down with policies
5. **Communicate relentlessly** — migration is as much about people as technology

Summary

This notebook contains **112 best practices** across 9 migration steps. The breakdown by priority:

Priority	Count	Guidance
Critical	81	Must implement. Skipping any of these risks data loss, access failures, or migration rollback.
Recommended	27	Implement unless there is a documented reason not to. These reduce risk and cost.
Optional	4	Implement if resources allow. These improve efficiency but are not blockers.

Use the 9-step framework table and the 11-step Order of Operations table at the top of this notebook to track progress and govern execution sequence. Each step's best practices align to the corresponding S2S notebook (S2S-01 through S2S-10) for full context and worked examples.

This notebook was AI-generated from community-submitted and publicly available sources. This notebook series is not officially supported by Dynatrace. Always verify information against official Dynatrace documentation.